

# INTRODUCTION TO APACHE SPARK

---

# Agenda

- Introduction to Spark
- Transformations and Actions
- Spark Architecture
- Hadoop and Spark
- Spark vs MapReduce

# References

- Apache Spark Site - <http://spark.apache.org/>
- Spark Mailing List
- Blogs
  - Cloudera – <http://blog.cloudera.com>
  - DataBricks – <http://databricks.com/blog>
  - MapR - <http://www.mapr.com/blog>
  - HortonWorks - <http://hortonworks.com/blog>

**BigData**

**Computation**

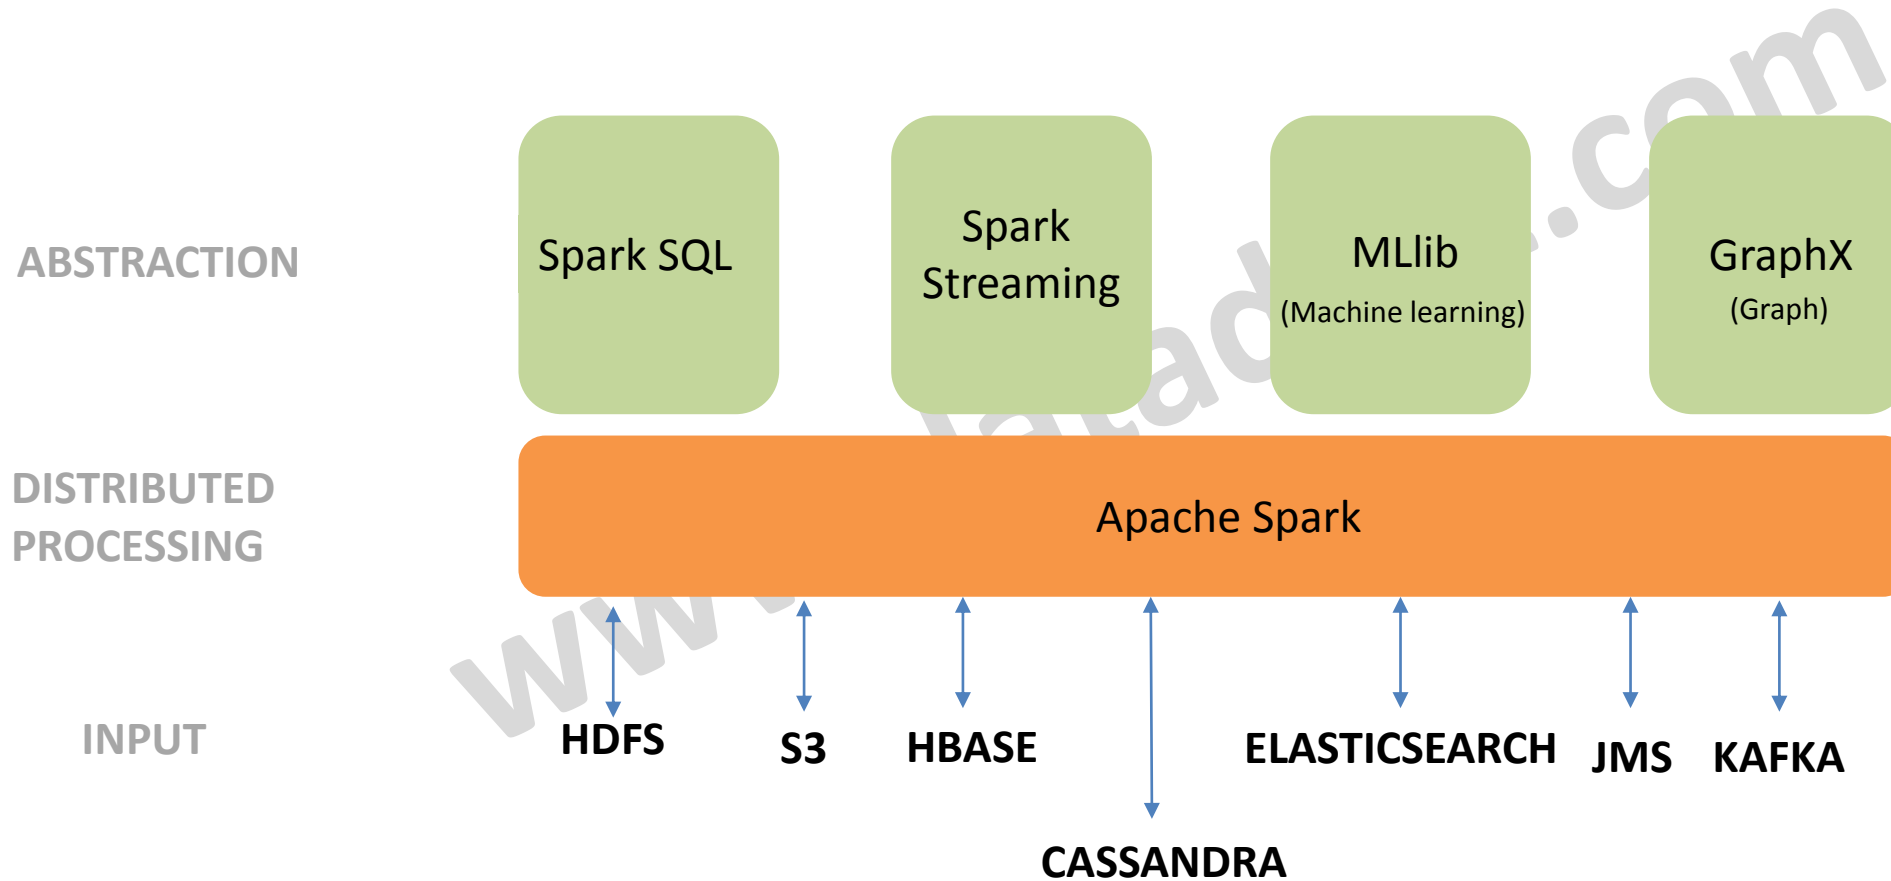
**Storage**

# What is Apache Spark

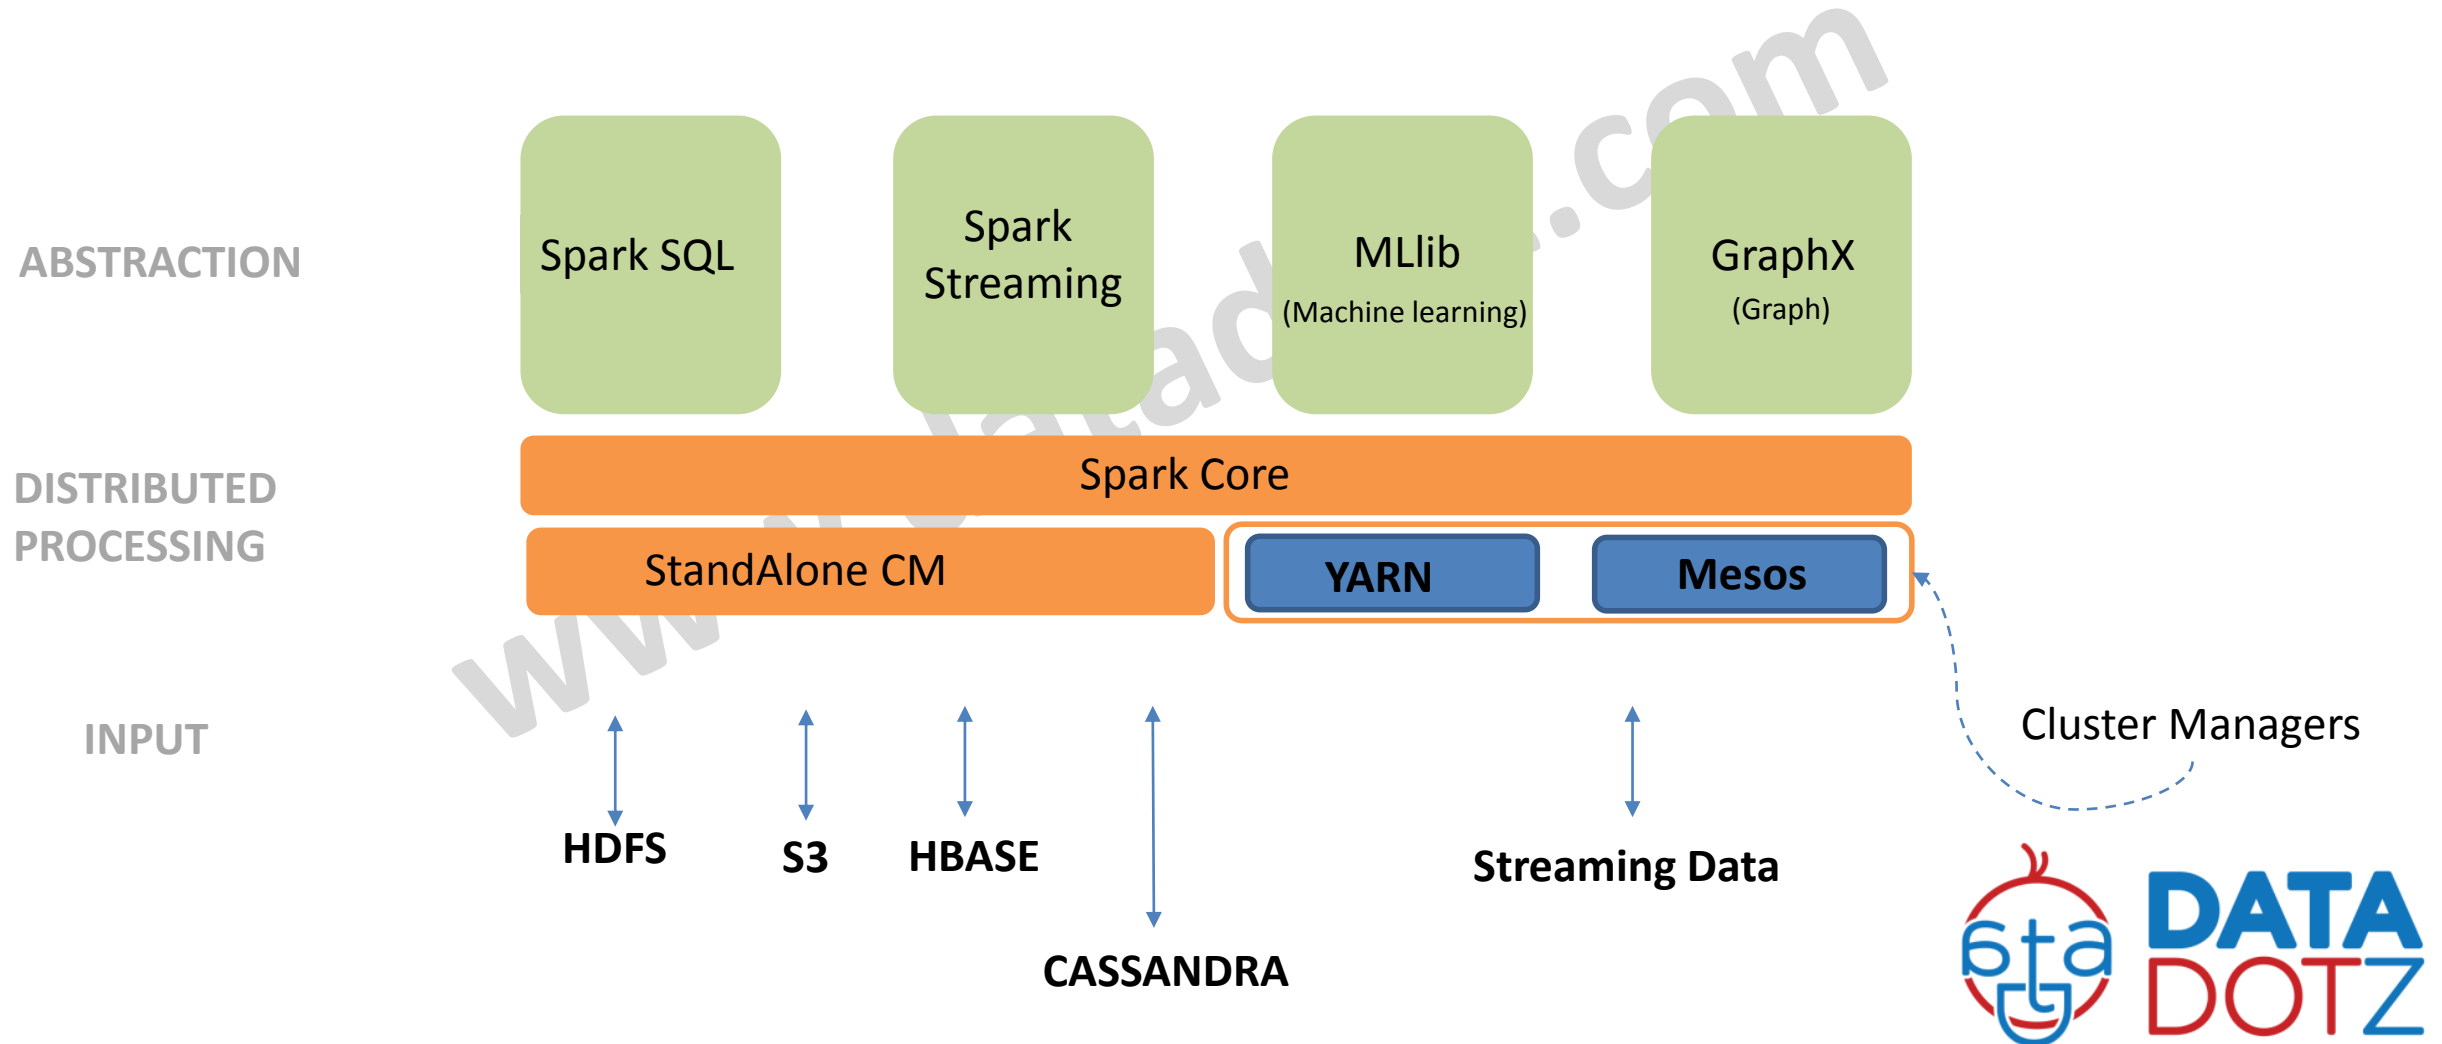
A Framework written in Scala which provides distributed data processing



# Spark Framework



# Spark Framework



# Spark officially sets a new record in large-scale sorting

	Hadoop MR Record	Spark Record	Spark 1 PB
Data Size	102.5 TB	100 TB	1000 TB
Elapsed Time	72 mins	23 mins	234 mins
# Nodes	2100	206	190
# Cores	50400 physical	6592 virtualized	6080 virtualized
Cluster disk throughput	3150 GB/s (est.)	618 GB/s	570 GB/s
Sort Benchmark Daytona Rules	Yes	Yes	No
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network	virtualized (EC2) 10Gbps network
Sort rate	1.42 TB/min	4.27 TB/min	4.27 TB/min
Sort rate/node	0.67 GB/min	20.7 GB/min	22.5 GB/min



Reference : <https://databricks.com/blog/2014/11/05/spark-officially-sets-a-new-record-in-large-scale-sorting.html>





# Who Uses Apache Spark?

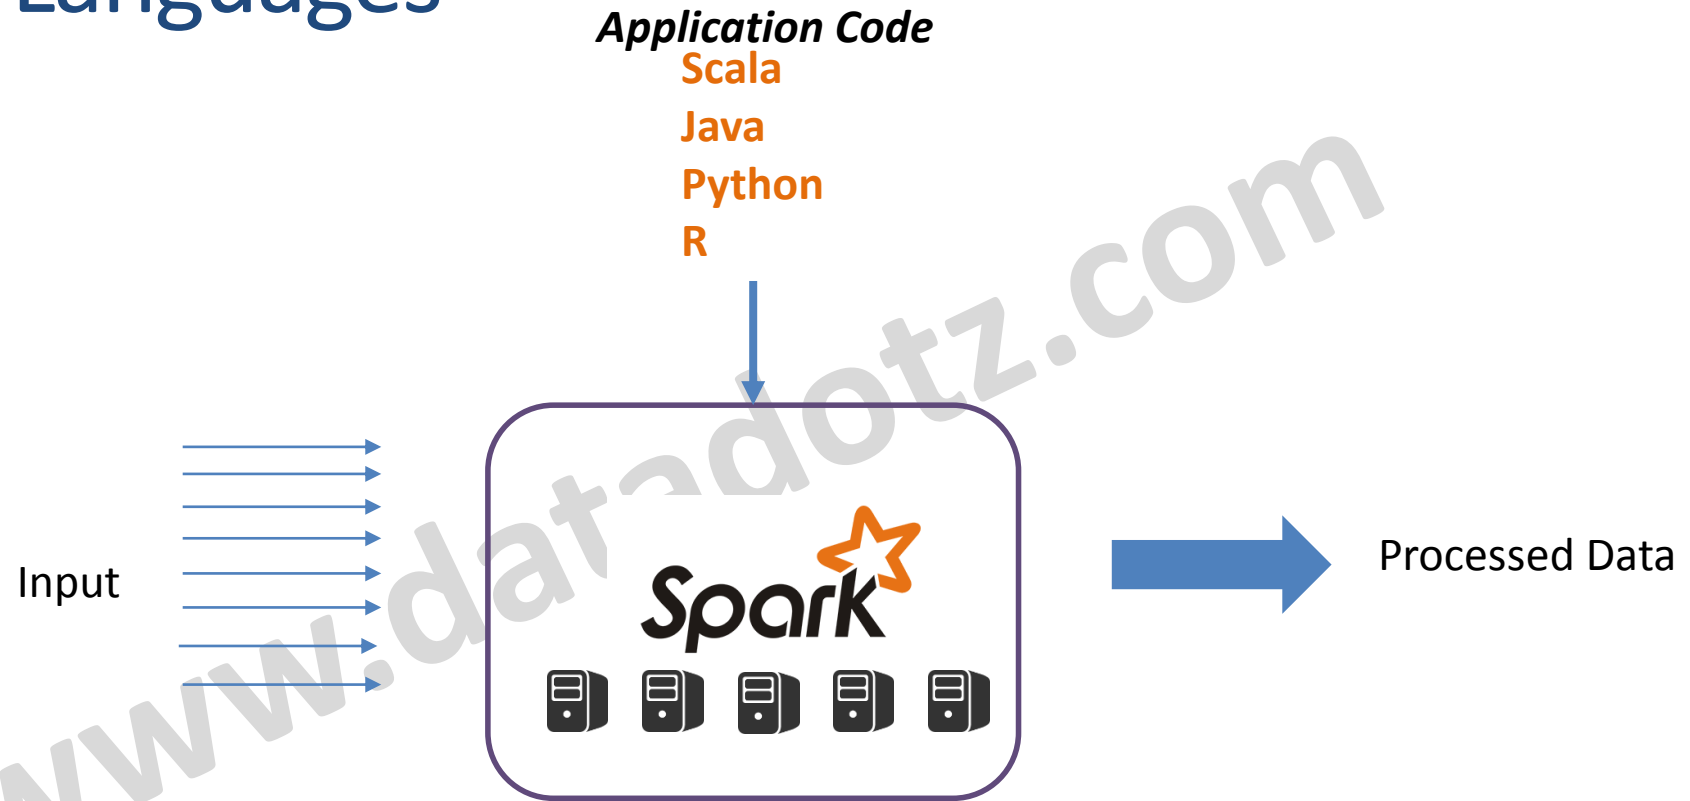


# History

- Started in UC Berkeley in 2009
  - Founder – Matei Zaharia
- Open Sourced in 2010 under BSD License
- 2011 – Higher Components as part of BDAS Stack
- 2013 – Joined ASF under Apache 2.0 License
- 2013 – DataBricks Founded
- Feb 2014 - Top Level Project in Apache
- May 2014 – Spark -1.0
- Nov 2014 - World record in large scale sorting by DataBricks Team



# Supported Languages



\*\*Spark Source Code is written in Scala

# Requirements

- OS
  - Windows, Linux, Mac OS
- JAVA\_HOME
- SCALA PATH

www.datadotz.com

# Installation Mode

- Interactive Shell for adhoc analysis or learning
  - Spark Shell – Interactive REPL
  - Can run locally or connect to a Spark cluster
- Cluster
  - Standalone
    - Amazon EC2
  - Mesos
  - YARN (Hadoop)

# Spark Installation

- Download
  - <https://spark.apache.org/downloads.html>
- Please build binary for your requirement.
  - For shell
    - *bin/sbt assembly*
    - *bin/sbt -Pyarn -Phadoop-2.6 -Phive -Phive-thriftserver assembly*
- Source Code
  - <http://github.com/apache/spark>

# Running Spark Shell

- Command
  - Scala - **bin/spark-shell**
  - Python – **bin/pyspark**

www.datadotz.com

# Spark Context

- Single Entry Point of the Spark Application (driver)
- Spark Shell provides a preconfigured Spark Context “**sc**”

## Scala


```
scala> sc.appName  
res1: String = Spark shell
```

## Python

```
>>> sc.appName  
u'PySparkShell'
```



```
[root@centovml spark-1.5.2]$ bin/spark-shell
log4j:WARN No appenders could be found for logger (org.apache.hadoop.metrics2.lib.MutableMetricsFactory).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Using Spark's repl log4j profile: org/apache/spark/log4j-defaults-repl.properties
To adjust logging level use sc.setLogLevel("INFO")
Welcome to
```



```
Using Scala version 2.10.4 (Java HotSpot(TM) 64-Bit Server VM, Java 1.7.0_67)
Type in expressions to have them evaluated.
Type :help for more information.
15/11/27 14:10:24 WARN MetricsSystem: Using default name DAGScheduler for source because spark.app.id is not set.
Spark context available as sc.
15/11/27 14:10:36 WARN ObjectStore: Version information not found in metastore. hive.metastore.schema.validation is not enabled so recording the schema version 1.2.0
15/11/27 14:10:36 WARN ObjectStore: Failed to get database default, returning NoSuchObjectException
15/11/27 14:10:39 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
SQL context available as sqlContext.

scala> █
```

← → ↻ 192.168.1.5:4040/jobs/ ☆ ABP ≡

**Spark** 1.5.2 Jobs Stages Storage Environment Executors SQL Spark shell application UI

## Spark Jobs (?)

Total Uptime: 49 s

Scheduling Mode: FIFO

▶ [Event Timeline](#)

# First Spark Program using Scala

```
scala> val data = Array(1, 2, 3, 4, 5)
```

```
data: Array[Int] = Array(1, 2, 3, 4, 5)
```

```
# create RDD by parallelizing the collection
```

```
# Alternative reading it from distributed storage such as HDFS, NoSQL, ..etc
```

```
scala> val distData = sc.parallelize(data)
```

```
distData: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:14
```

Fault Tolerant  
(recover)

Across the machines  
(partitioned)

Collection of Records  
(Immutable)

## Resilient Distributed DataSets

- Read Only , partitioned collection of records



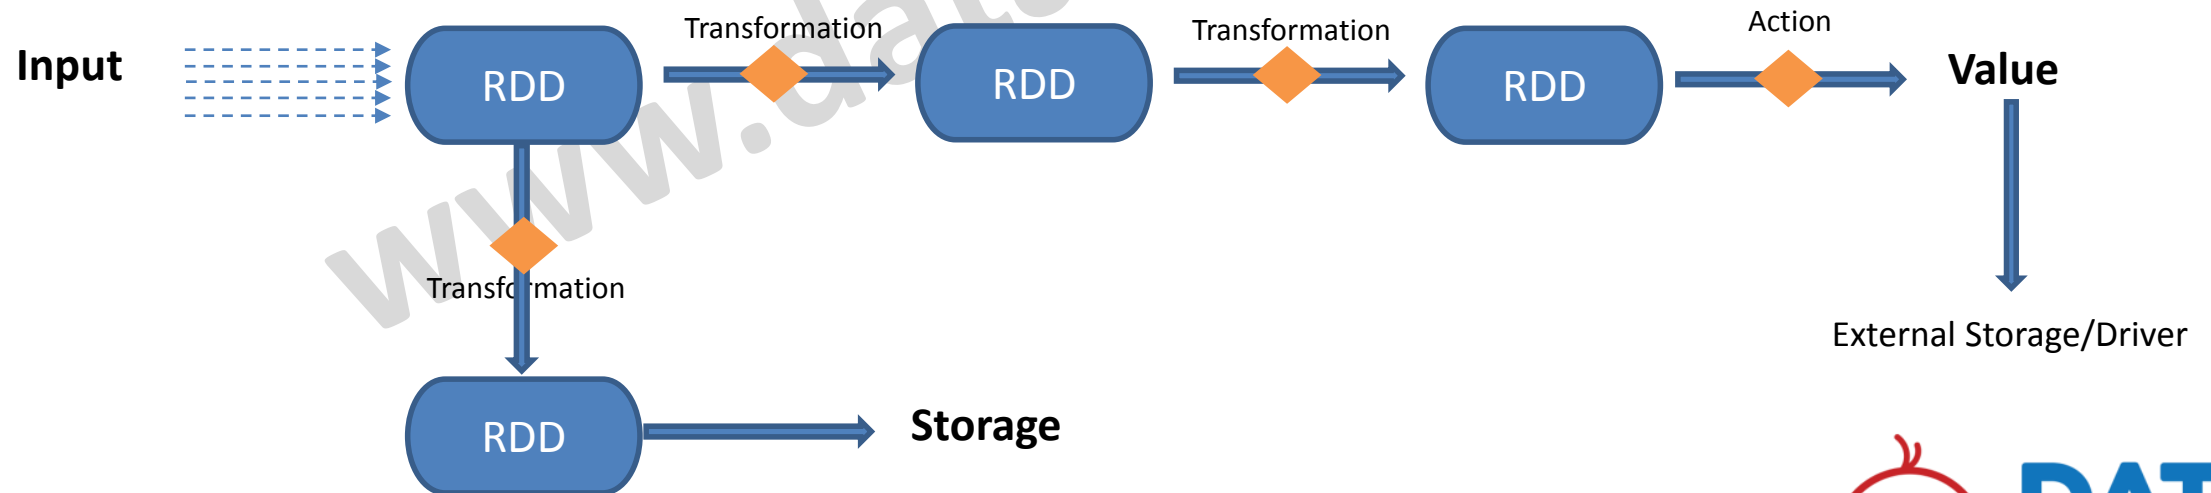
# RDD operations

## Transformations

- Create a New DataSet from a New DataSet
- Transformations are lazy operations
- RDDs are operated when an action is run on it.

## Actions

- Compute Values
- Return Values or write Output to external Storage
- Earlier Transformations are applied to RDD
  - Since transformations are lazy operations



# Select only records less than 3

```
scala> val data = Array(1, 2, 3, 4, 5)
```

```
data: Array[Int] = Array(1, 2, 3, 4, 5)
```

```
# create RDD
```

```
scala> val distData = sc.parallelize(data)
```

```
distData: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:23
```

```
# filter RDDs for elements less than 3
```

```
# Transformation – create RDDs from existing RDDs
```

```
# Transformation are Lazy operations
```

```
scala> val filteredData = distData.filter( _ < 3 )
```

```
filteredData: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[1] at filter at <console>:25
```

```
# collect or obtain output RDDs
```

```
# Actions – calculate values from RDDs
```

```
scala> val resultArray= filteredData.collect()
```

```
resultArray: Array[Int] = Array(1, 2)
```



# Lineage of RDDs – Fault Recovery

- RDDs maintain their Lineage
  - Any intermediate RDDs is missing , it can calculate from its Parent RDDs

```
scala> val data = Array(1, 2, 3, 4, 5)
```

```
data: Array[Int] = Array(1, 2, 3, 4, 5)
```

```
scala> val distData = sc.parallelize(data)
```

```
distData: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:23
```

```
scala> val filteredData = distData.filter(_ < 3)
```


```
filteredData: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[1] at filter at <console>:25
```

```
scala> filteredData.toDebugString
```

```
res3: String =
```

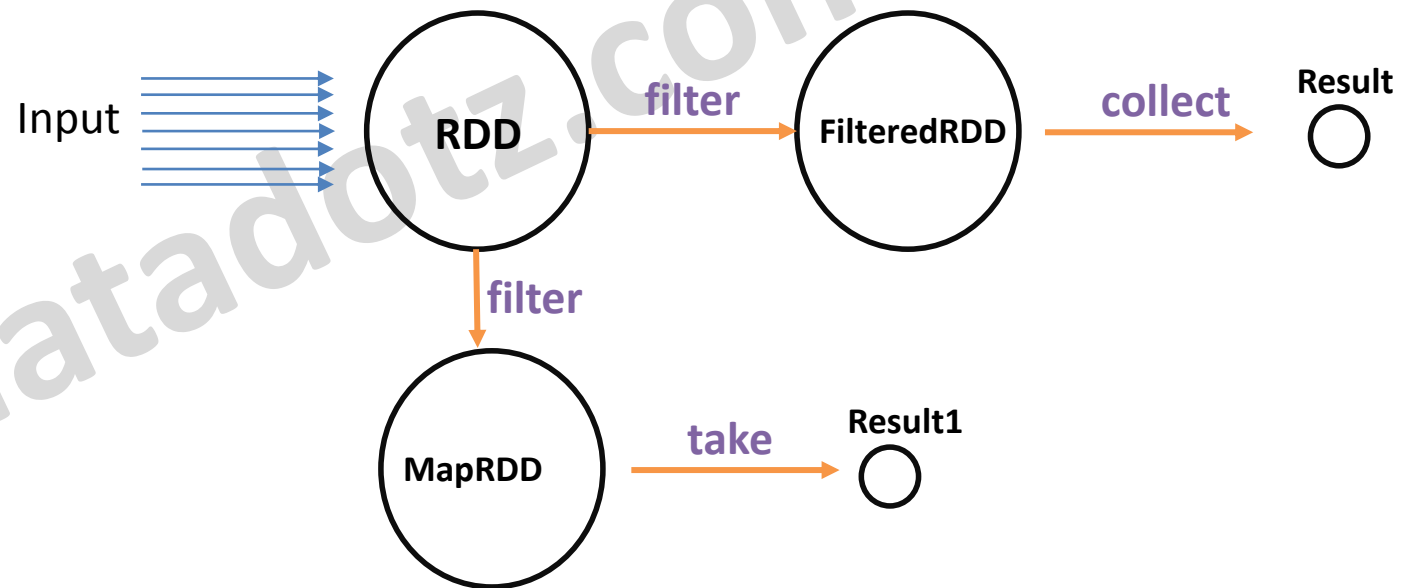
```
(1) MapPartitionsRDD[2] at filter at <console>:25 []
```

```
| ParallelCollectionRDD[0] at parallelize at <console>:23 []
```



# Directed Acyclic Graph of RDDs

- Directed
  - Only in a Single Direction
- Acyclic
  - No Looping
- Provides Fault-Tolerance
  - By providing **Lineage**



## Filter – contd..

```
scala> val distData = sc.parallelize(Array(1, 2, 3, 4, 5))
```

```
distData: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:23
```

# filter – takes a condition function

# Any RDD which satisfies the condition will added to resultant RDD

```
scala> val filteredData = distData.filter( i => i %2 == 0)
```

```
filteredData: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[1] at filter at <console>:25
```

```
scala> filteredData.collect()
```

```
res1: Array[Int] = Array(2, 4)
```



# Partitions

```
scala> val distData = sc.parallelize(Array(1, 2, 3, 4, 5))
```

```
distData: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:23
```

```
# RDD -> read only partitioned collection of records
```

```
scala> distData.partitions.length
```

```
res8: Int = 1
```

```
# custom parallelism with custom partitions
```

```
scala> val distData = sc.parallelize(Array(1, 2, 3, 4, 5), 3)
```

```
distData: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:23
```

```
scala> distData.partitions.length
```

```
res9: Int = 3
```

# Data From Local file

```
1,Brandon Buckner,avil,female,525
2,Veda Hopkins,avil,male,633
3,Zia Underwood,paracetamol,male,980
4,Austin Mayer,paracetamol,female,338
5,Mara Higgins,avil,female,153
6,Sybill Crosby,avil,male,193
7,Tyler Rosales,paracetamol,male,778
8,Ivan Hale,avil,female,454
9,Alika Gilmore,paracetamol,female,833
10,Len Burgess,metacin,male,325
```

Select drug, sum(amount) from patient group by drug;

```
scala> val patient = sc.textFile("/home/user/data/patient.txt")
```

```
patient: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[11] at textFile at <console>:21
```

```
scala> val mappedKVs = patient.map(line => (line.split(",")(2),line.split(",")(4).toInt))
```

```
mappedKVs: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[13] at map at <console>:23
```

```
scala> val result = mappedKVs.reduceByKey(_+_ , 1)
```

```
result: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[14] at reduceByKey at <console>:25
```

```
scala> result.collect()
```

```
res2: Array[(String, Int)] = Array((avil,1958), (metacin,325), (paracetamol,2929))
```



# Basic Actions in Spark

# Actions

```
scala> val distData = sc.parallelize(Array(1, 2, 3, 4, 5))
```

```
scala> distData.collect()  
res3: Array[Int] = Array(1, 2, 3, 4, 5)
```

```
scala> distData.take(3)  
res0: Array[Int] = Array(1, 2, 3)
```

```
scala> distData.take(3)  
res0: Array[Int] = Array(1, 2, 3)
```

```
scala> distData.top(1)  
res1: Array[Int] = Array(5)
```

```
scala> distData.reduce(_+_)  
res2: Int = 15
```

```
scala> distData.first()  
res0: Int = 1
```

```
scala> distData.count()  
res2: Long = 5
```



# Actions

```
scala> val distData = sc.parallelize(Array(1, 2, 3, 4, 5))
```

```
scala> println("Hello, world!")  
Hello, world!
```

```
scala> distData.foreach(println)
```

```
11  
12  
13  
14  
15
```

```
scala> distData.saveAsTextFile("/home/user/resultdir")
```

# Actions

reduce(function)

collect()

count()

countByValue()

first()

take(n)

takeSample(withReplacement, num, [seed])

takeOrdered(n, [ordering])

saveAsTextFile(path)

saveAsSequenceFile(path)

saveAsObjectFile(path)

countByKey()

foreach(function)

foreachPartition()

treeAggregate()

treeReduce()

# Basic Transformations in Spark

# Transformation -

```
scala> val distData = sc.parallelize(Array(1, 2, 3, 4, 5))
```

```
distData: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[2] at parallelize at <console>:12
```

```
scala> val mappedData = distData.map(x => x*x).collect()
```

```
res1: Array[Int] = Array(1, 4, 9, 16, 25)
```

```
scala> val filteredData = distData.filter(x => x>= 3).collect()
```

```
res1: Array[Int] = Array(3, 4, 5)
```

```
scala> distData.flatMap(_>_.toUpperCase).collect()
```

```
res4: Array[Char] = Array(A, P, P, L, E, T, E, S, T)
```



# KeyValue Pair

## Scala

```
scala> val kvData = (1, "senthil")
```

```
kvData: (Int, String) = (1,senthil)
```

```
scala> kvData._1
```

```
res2: Int = 1
```

```
scala> kvData._2
```

```
res3: String = senthil
```

## Python

```
>>> kvData = (1,"senthil")
```

```
>>> kvData[0]
```

```
1
```

```
>>> kvData[1]
```

```
'senthil'
```

## Java

```
import scala.Tuple2;
```

```
.
```

```
.
```

```
Tuple2 kvData = new Tuple2(1,"senthil");
```

```
System.out.println(kvData._1);
```

```
System.out.println(kvData._2);
```

Tuple - sequence of immutable objects

In Scala – Tuple2 forms KeyValue Pair in Spark



# Basic Key Value Pairs

```
scala> val Data= sc.parallelize(Array("Apple-1","orange-3","Apple-4"))
```

```
Data: org.apache.spark.rdd.RDD[String] = ParallelCollectionRDD[14] at parallelize at <console>:12
```

```
# create Key Value pair RDD
```

```
# map (Transformation) – converts RDD into KeyValue Pair
```

```
scala> val mappedKVs = Data.map(element => (element.split("-")(0), element.split("-")(1).toInt))
```

```
mappedKVs: org.apache.spark.rdd.RDD[(String, Int)] = MappedRDD[15] at map at <console>:14
```

```
# collect all KV pair RDDs
```

```
scala> mappedKVs.collect()
```

```
res12: Array[(String, Int)] = Array((Apple,1), (orange,3), (Apple,4))
```

```
# collect Keys alone
```

```
scala> mappedKVs.keys.collect()
```

```
res8: Array[String] = Array(Apple, orange, Apple)
```

```
# collect Values alone
```

```
scala> mappedKVs.values.collect()
```

```
res9: Array[Int] = Array(1, 3, 4)
```



# Basic Key/Value Transformations

```
scala> val kvData= sc.parallelize(Array(("Apple",1),("orange",3),("Apple",4)))
```

```
kvData: org.apache.spark.rdd.RDD[(String, Int)] = ParallelCollectionRDD[8] at parallelize at <console>:12
```

```
scala> kvData.reduceByKey(_+_).collect()
```

```
res7: Array[(String, Int)] = Array((orange,3), (Apple,5))
```

```
scala> kvData.groupByKey().collect()
```

```
res8: Array[(String, Iterable[Int])] = Array((orange,CompactBuffer(3)), (Apple,CompactBuffer(1, 4)))
```

```
scala> kvData.sortByKey().collect()
```

```
res9: Array[(String, Int)] = Array((Apple,1), (Apple,4), (orange,3))
```

```
scala> kvData.sortByKey(false).collect()
```

```
res0: Array[(String, Int)] = Array((orange,3), (Apple,1), (Apple,4))
```

# Other Transformations on PairRDD

- subtractByKey
  - Remove elements for keys in second RDD
- Joins
  - Works on two RDDS
  - Join, RightOuterJoin, LeftOuterJoin
  - Internally cogroup. Cogroup can be used to work on **more than** two RDDs at the same time.

# Additional Actions on PairRDD

```
scala> val kvData= sc.parallelize(Array(("Apple",1),("orange",3),("Apple",4)))
```

```
kvData: org.apache.spark.rdd.RDD[(String, Int)] = ParallelCollectionRDD[8] at parallelize at <console>:12
```

```
scala> kvData.countByKey()
```

```
res9: scala.collection.Map[String,Long] = Map(orange -> 1, Apple -> 2)
```

```
scala> kvData.collectAsMap()
```

```
res10: scala.collection.Map[String,Int] = Map(orange -> 3, Apple -> 4)
```

```
scala> kvData.lookup("Apple")
```

```
res11: Seq[Int] = WrappedArray(1, 4)
```

# Transformations

map(*function*)

filter(*function*)

filterByRange(*lower, upper*)

flatMap(*function*)

mapPartitions(*function*)

mapPartitionsWithIndex(*function*)

sample(*withReplacement, fraction, seed*)

union(*otherDataset*)

intersection(*otherDataset*)

distinct(*[numTasks]*)

groupByKey(*[numTasks]*)

reduceByKey(*function, [numTasks]*)

aggregateByKey(*zeroValue*)(*seqOp, combOp, [numTasks]*)

sortByKey(*[ascending], [numTasks]*)

join(*otherDataset, [numTasks]*)

cogroup(*otherDataset, [numTasks]*)

cartesian(*otherDataset*)

pipe(*command, [envVars]*)

coalesce(*numPartitions*)

repartition(*numPartitions*)

repartitionAndSortWithinPartitions(*partitioner*)

# Numerical RDD Operations

```
scala> val distData = sc.parallelize(Array(1, 2, 3, 4, 5))
```

```
kvData: org.apache.spark.rdd.RDD[(String, Int)] = ParallelCollectionRDD[8] at parallelize at <console>:12
```

```
# Statistics Operations on the Data
```

```
# Returned StatsCounter object by calling stats
```

```
scala> distData.stats()
```

```
res7: org.apache.spark.util.StatCounter = (count: 5, mean: 3.000000, stdev: 1.414214, max: 5.000000, min: 1.000000)
```

```
# Call direct methods if needed
```

```
scala> distData.max()
```

```
res8: Int = 5
```

# Logical DAG

```
scala> val Data= sc.parallelize(Array("Apple-1","orange-3","Apple-4"))  
Data: org.apache.spark.rdd.RDD[String] = ParallelCollectionRDD[9] at parallelize at <console>:21
```

```
scala> val mappedKVs = Data.map(line => (line.split("-")(0),line.split("-")(1).toInt))  
mappedKVs: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[7] at map at <console>:23
```

```
scala> val result = mappedKVs.reduceByKey(_+_)  
result: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[9] at reduceByKey at <console>:25
```

```
scala> result.toDebugString
```

```
res4: String =
```

```
(1) ShuffledRDD[11] at reduceByKey at <console>:25 []
```

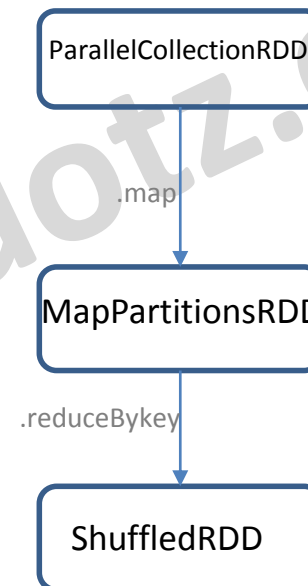
```
+- (1) MapPartitionsRDD[[10] at map at <console>:23 []
```

```
    | ParallelCollectionRDD[9] at parallelize at <console>:21 []
```

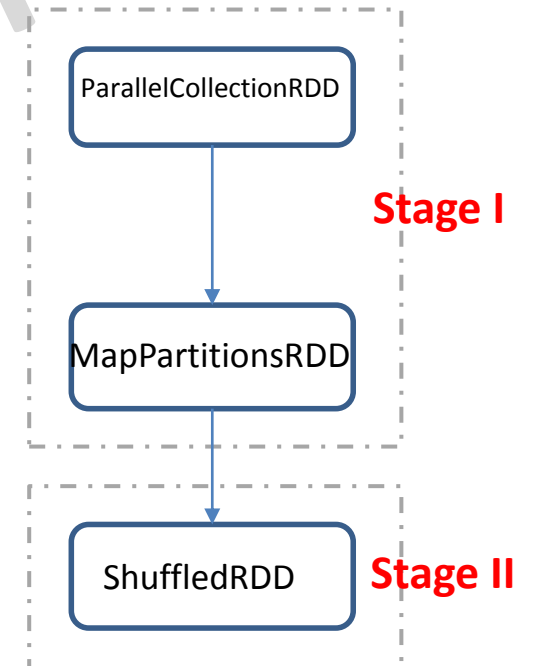
Stage II

Stage I

## DAG Graph



## Physical Plan





# Application -> Jobs -> Stages -> Tasks

- Application contains only one Spark Context
- For Every Action, it creates a Job
- Each Job consists of Stages
- Stage consists of Tasks for each partition in that RDD
  - Tasks -> computation on each partition of the Data

```
scala> val Data= sc.parallelize(Array("Apple-1","orange-3","Apple-4"))
Data: org.apache.spark.rdd.RDD[String] = ParallelCollectionRDD[9] at parallelize at <console>:21
```

```
scala> val mappedKVs = Data.map(line => (line.split("-")(0),line.split("-")(1).toInt))
mappedKVs: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[7] at map at <console>:23
```

```
scala> val result = mappedKVs.reduceByKey(_+_)
result: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[9] at reduceByKey at <console>:25
```

```
scala> result.toDebugString
```

```
res4: String =
```

```
(1) ShuffledRDD[11] at reduceByKey at <console>:25 []
```

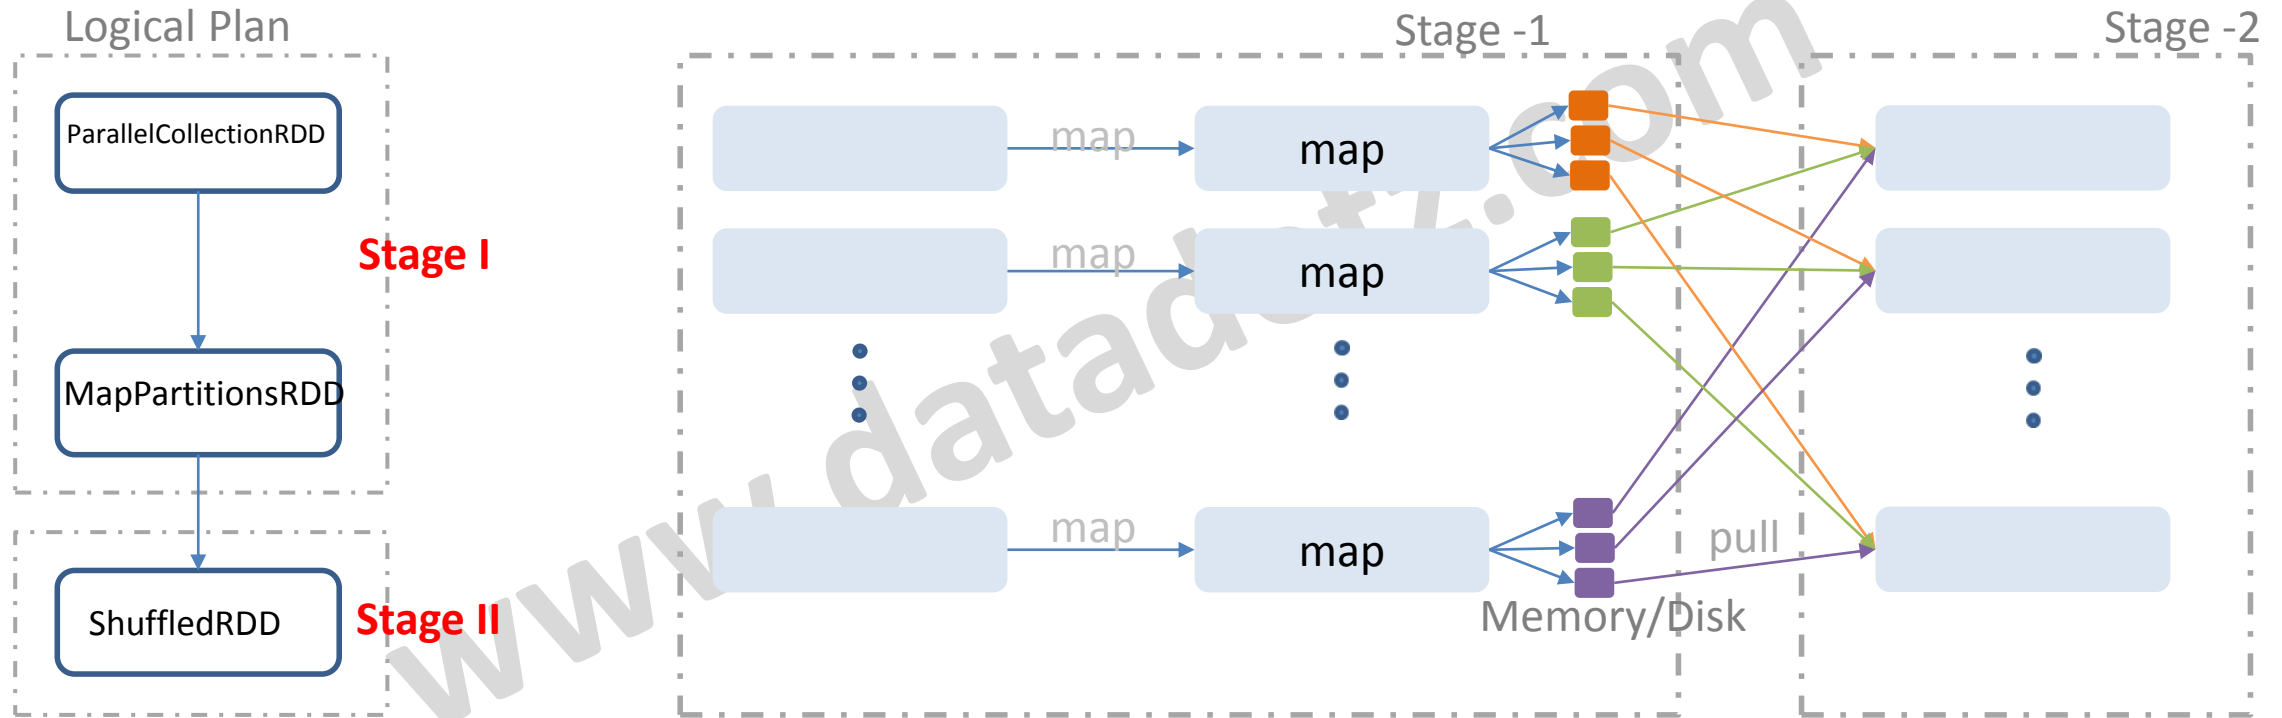
```
+-(1) MapPartitionsRDD[[10] at map at <console>:23 []
```

```
| ParallelCollectionRDD[9] at parallelize at <console>:21 []
```

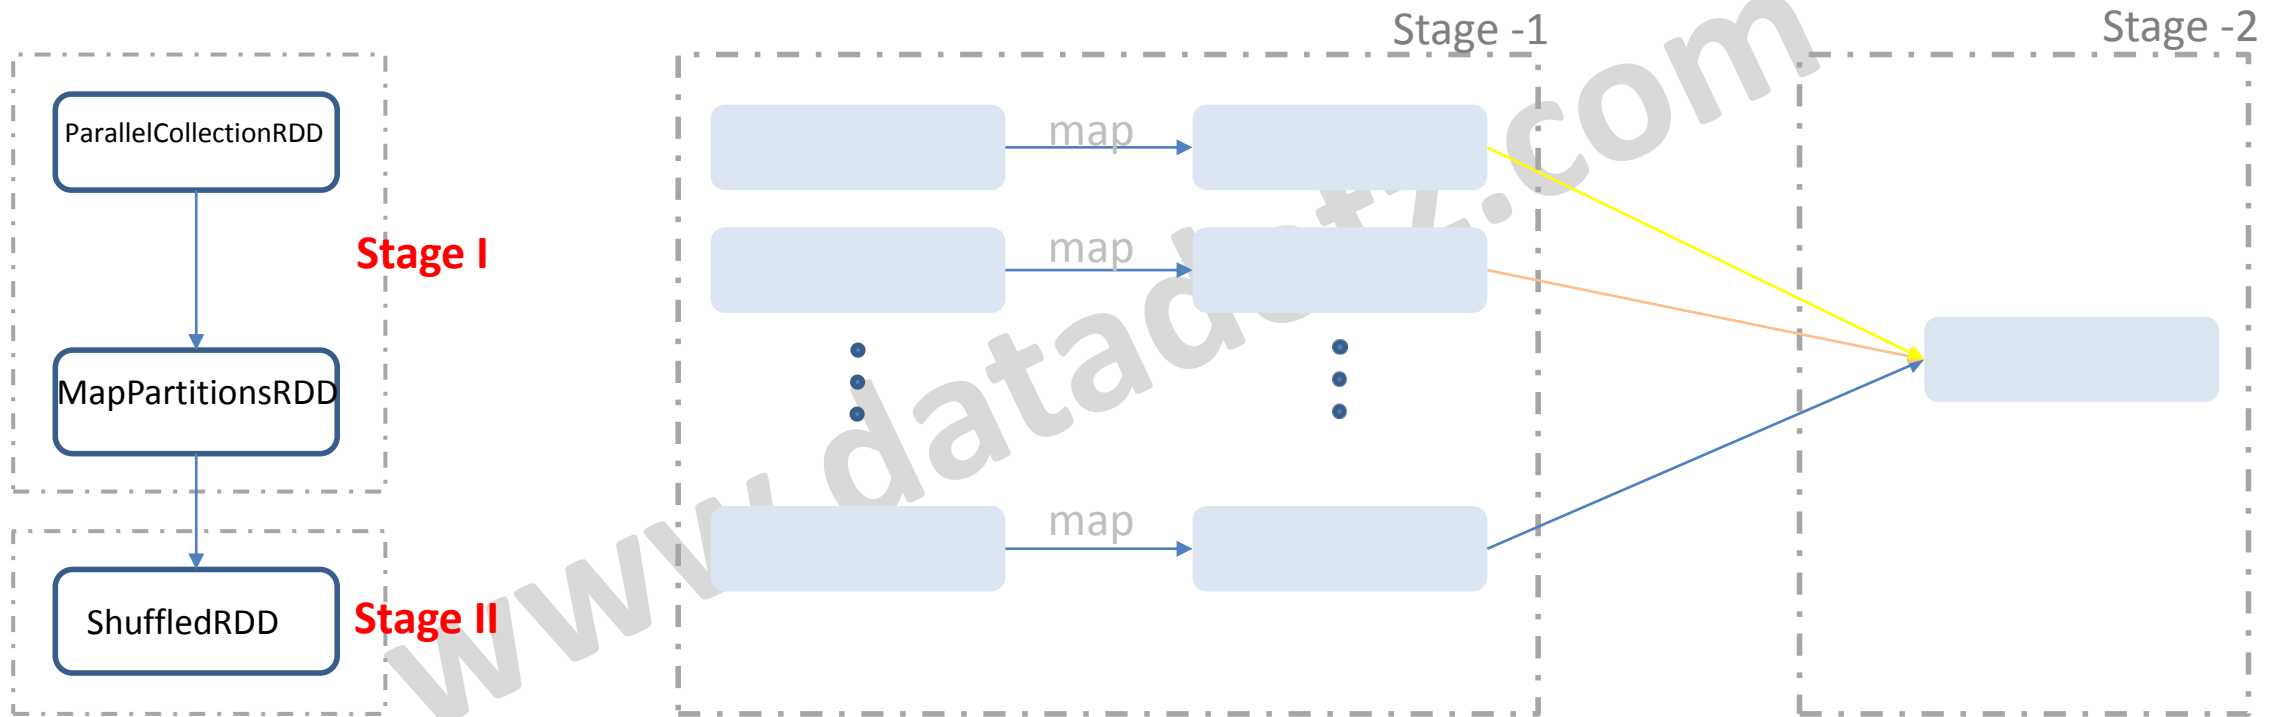
Stage II

Stage I

# Typical Physical Plan



# Typical Physical Plan



`mappedKVs.reduceByKey(_+_ ,1)`

*Task = data Partition + Computation*

# Data From Local file

```
1,Brandon Buckner,avil,female,525
2,Veda Hopkins,avil,male,633
3,Zia Underwood,paracetamol,male,980
4,Austin Mayer,paracetamol,female,338
5,Mara Higgins,avil,female,153
6,Sybill Crosby,avil,male,193
7,Tyler Rosales,paracetamol,male,778
8,Ivan Hale,avil,female,454
9,Alika Gilmore,paracetamol,female,833
10,Len Burgess,metacin,male,325
```

Select drug, sum(amount) from patient group by drug;

```
scala> val patient = sc.textFile("/home/user/data/patient.txt")
```

```
patient: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[11] at textFile at <console>:21
```

```
scala> val mappedKVs = patient.map(line => (line.split(",")(2),line.split(",")(4).toInt))
```

```
mappedKVs: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[13] at map at <console>:23
```

```
scala> val result = mappedKVs.reduceByKey(_+_ , 1)
```

```
result: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[14] at reduceByKey at <console>:25
```

```
scala> result.collect()
```

```
res2: Array[(String, Int)] = Array((avil,1958), (metacin,325), (paracetamol,2929))
```

# Where Clause = filter in Spark

Select drug, sum(amount) from patient where drug == "avil" group by drug;

```
scala> val patient = sc.textFile("/home/user/data/patient.txt")
```

*patient: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[11] at textFile at <console>:21*

```
scala> val mappedKVs = patient.filter(_.split(",")(2) == "avil").map(line => (line.split(",")(2), line.split(",")(4).toInt))
```

*mappedKVs: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[13] at map at <console>:23*

```
scala> val result = mappedKVs.reduceByKey(_+_ , 1)
```

*result: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[14] at reduceByKey at <console>:25*

```
scala> result.collect()
```

*res2: Array[(String, Int)] = Array((avil,1958))*

# distinct

Select distinct(drug) from patient;

```
scala> val patient = sc.textFile("/home/user/data/patient.txt")
```

*patient: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[11] at textFile at <console>:21*

```
scala> val drug_distinct= patient.map(line => line.split(",")(2)).distinct
```

*drug\_distinct: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[11] at textFile at <console>:21*

```
scala> drug_distinct.foreach(println)
```

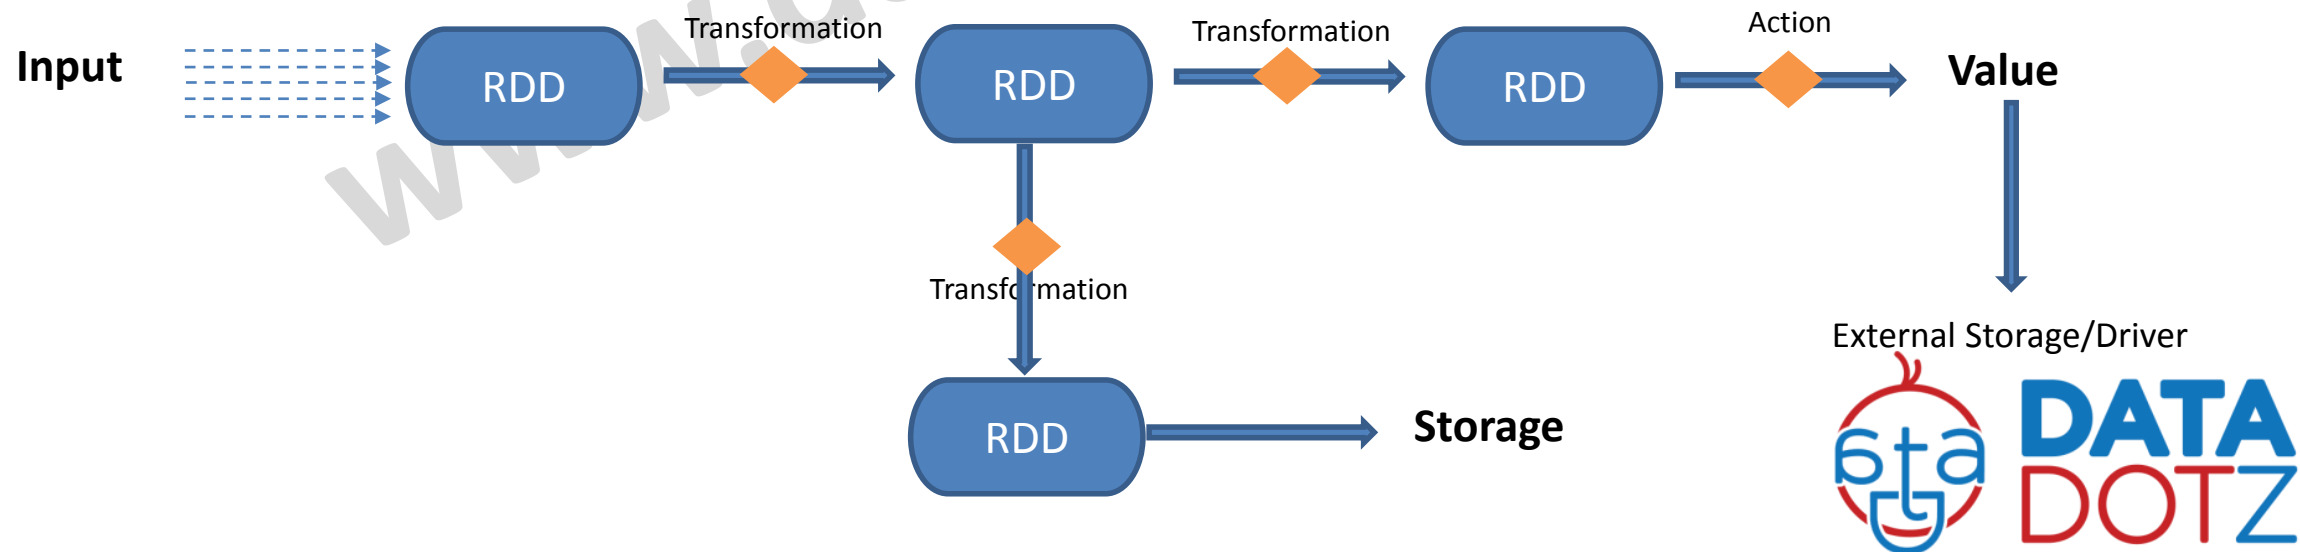
avil

metacin

paracetamol

# RDD Persistence /Caching

- Avoid re-evaluation of RDD, Spark provides many levels of Storing RDDs
  - MEMORY\_ONLY, MEMORY\_AND\_DISK , DISK\_ONLY
- To Persist/cache RDD, use below methods
  - persist()
  - cache() – use persist with MEMORY LEVEL
  - Both are lazy operations



Storage Level	Description	Format
MEMORY_ONLY (default)	<i>Recomputed if it does not fit in Memory</i>	
MEMORY_AND_DISK	<i>Spill to Disk on memory Full</i>	
MEMORY_ONLY_SER	<i>Recomputed if it does not fit in Memory</i>	<i>serialized</i>
MEMORY_AND_DISK_SER	<i>Spill to Disk on memory Full</i>	<i>serialized</i>
DISK_ONLY	<i>RDD Partitions in Disk</i>	
MEMORY_ONLY_2, MEMORY_AND_DISK_2		
OFF_HEAP (experimental)	<i>In <b>Tachyon</b> (distributed memory centric FileSystem)</i>	<i>serialized</i>



# Persistence APIs

cache()

persist([Storage Level])

unpersist()

checkpoint()\*\*

isCheckpointed()\*\*

getCheckpointFile()\*\*


# Removing Data from cache

- Removes old partitions of DataSet (RDDs) in LRU fashion per node basis Automatically - cache
- Remove RDDs manually by calling method
  - *RDD.unpersist()* method
  - Acts immediately

# Shared Variables

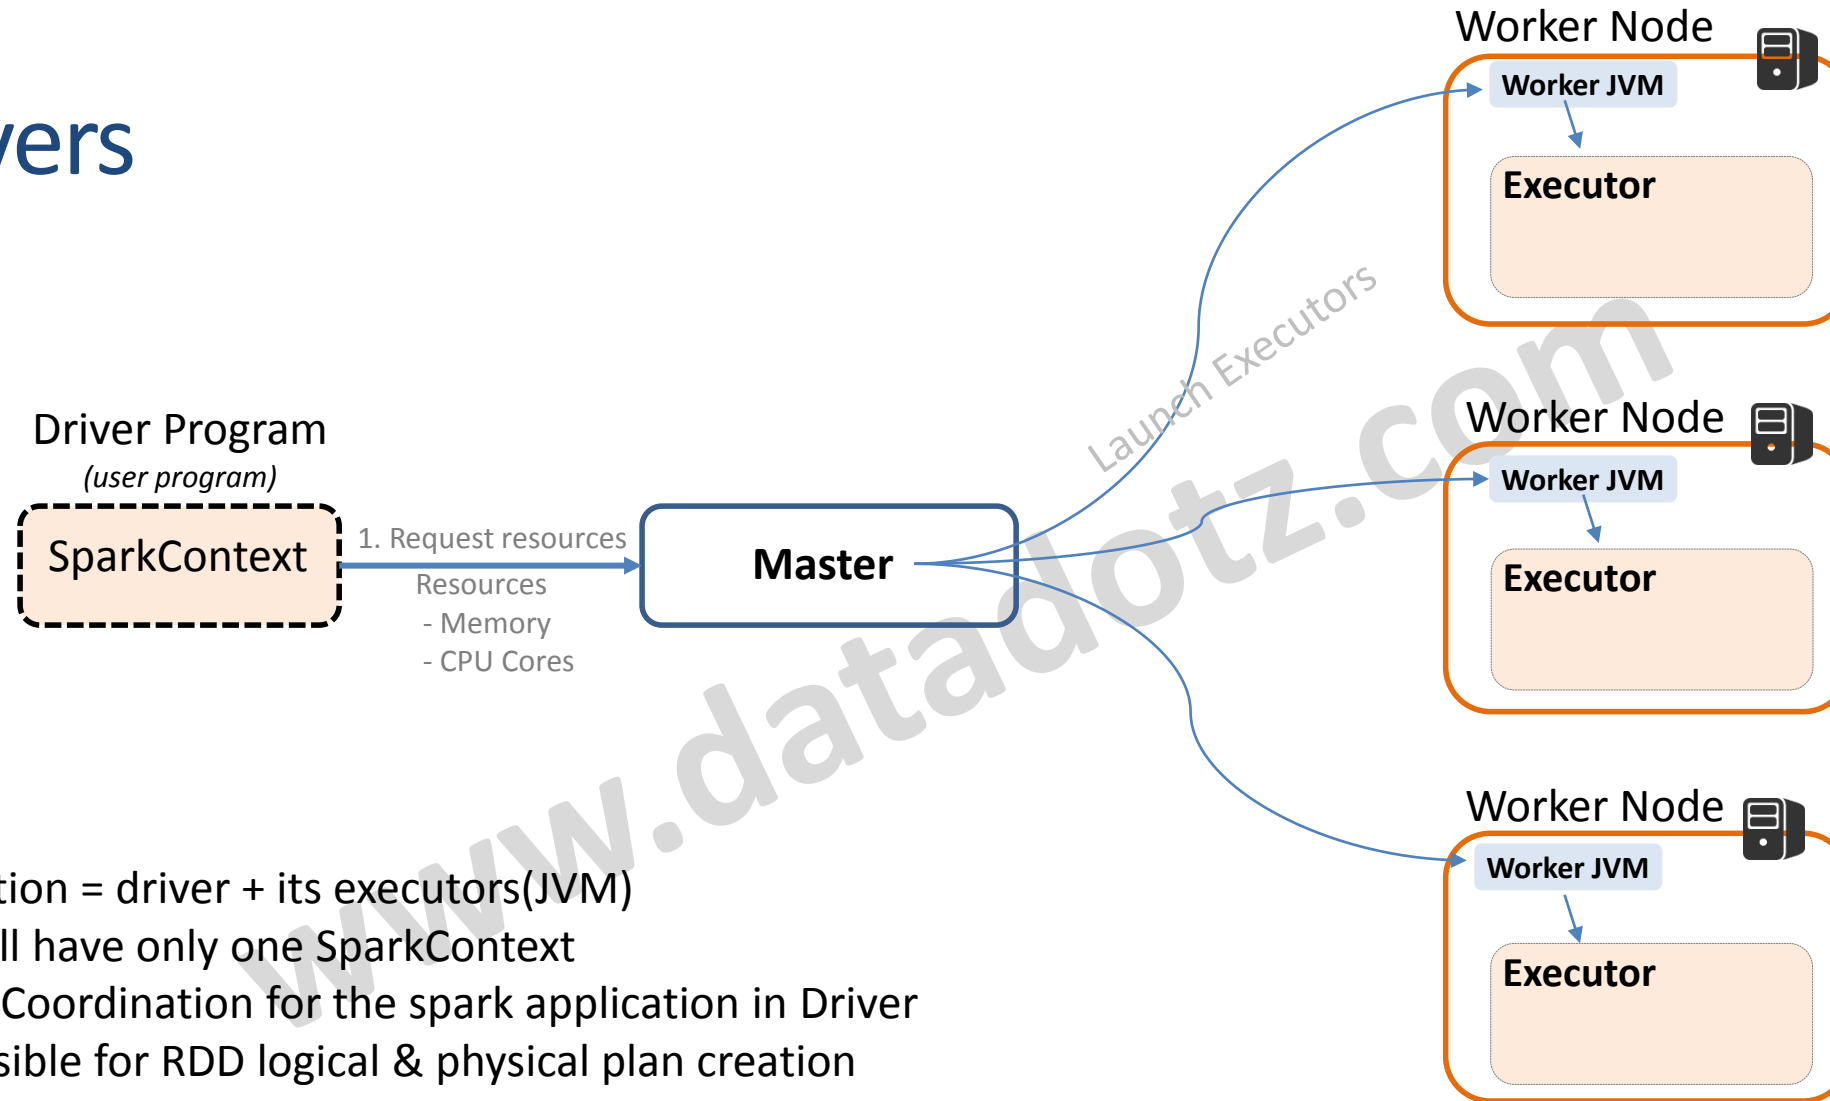
- Broadcast variables
  - To Share read only variables
- Accumulators
  - Aggregate information
  - Similar to Hadoop custom Counter
  - Accumulators can be seen in WebUI
  - Each task will have local accumulators
    - Spark will update each task's update to global accumulator only once if used in actions
    - If used in transformations, it may result in irregular values.

# StandAlone Spark Cluster



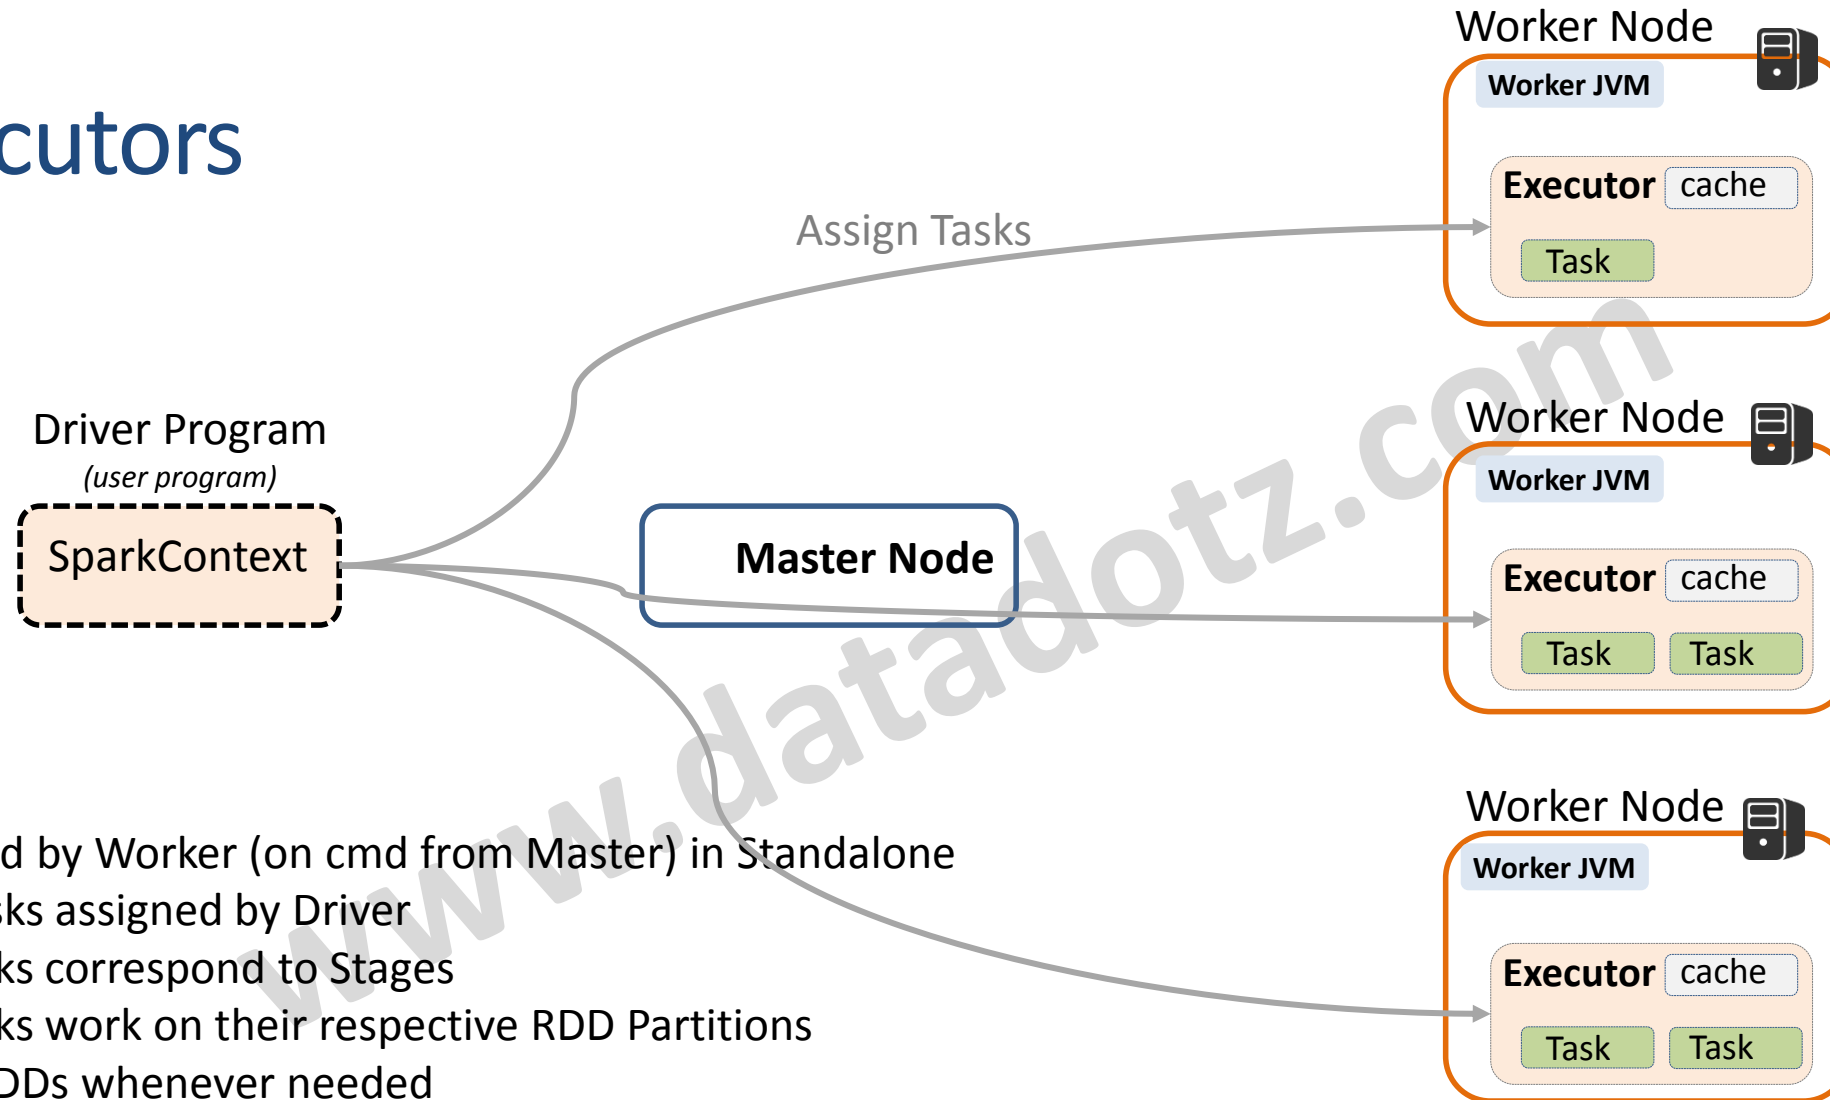
- Master / Slave Architecture
- Daemons(JVM)
  - Master
  - Worker

# Drivers



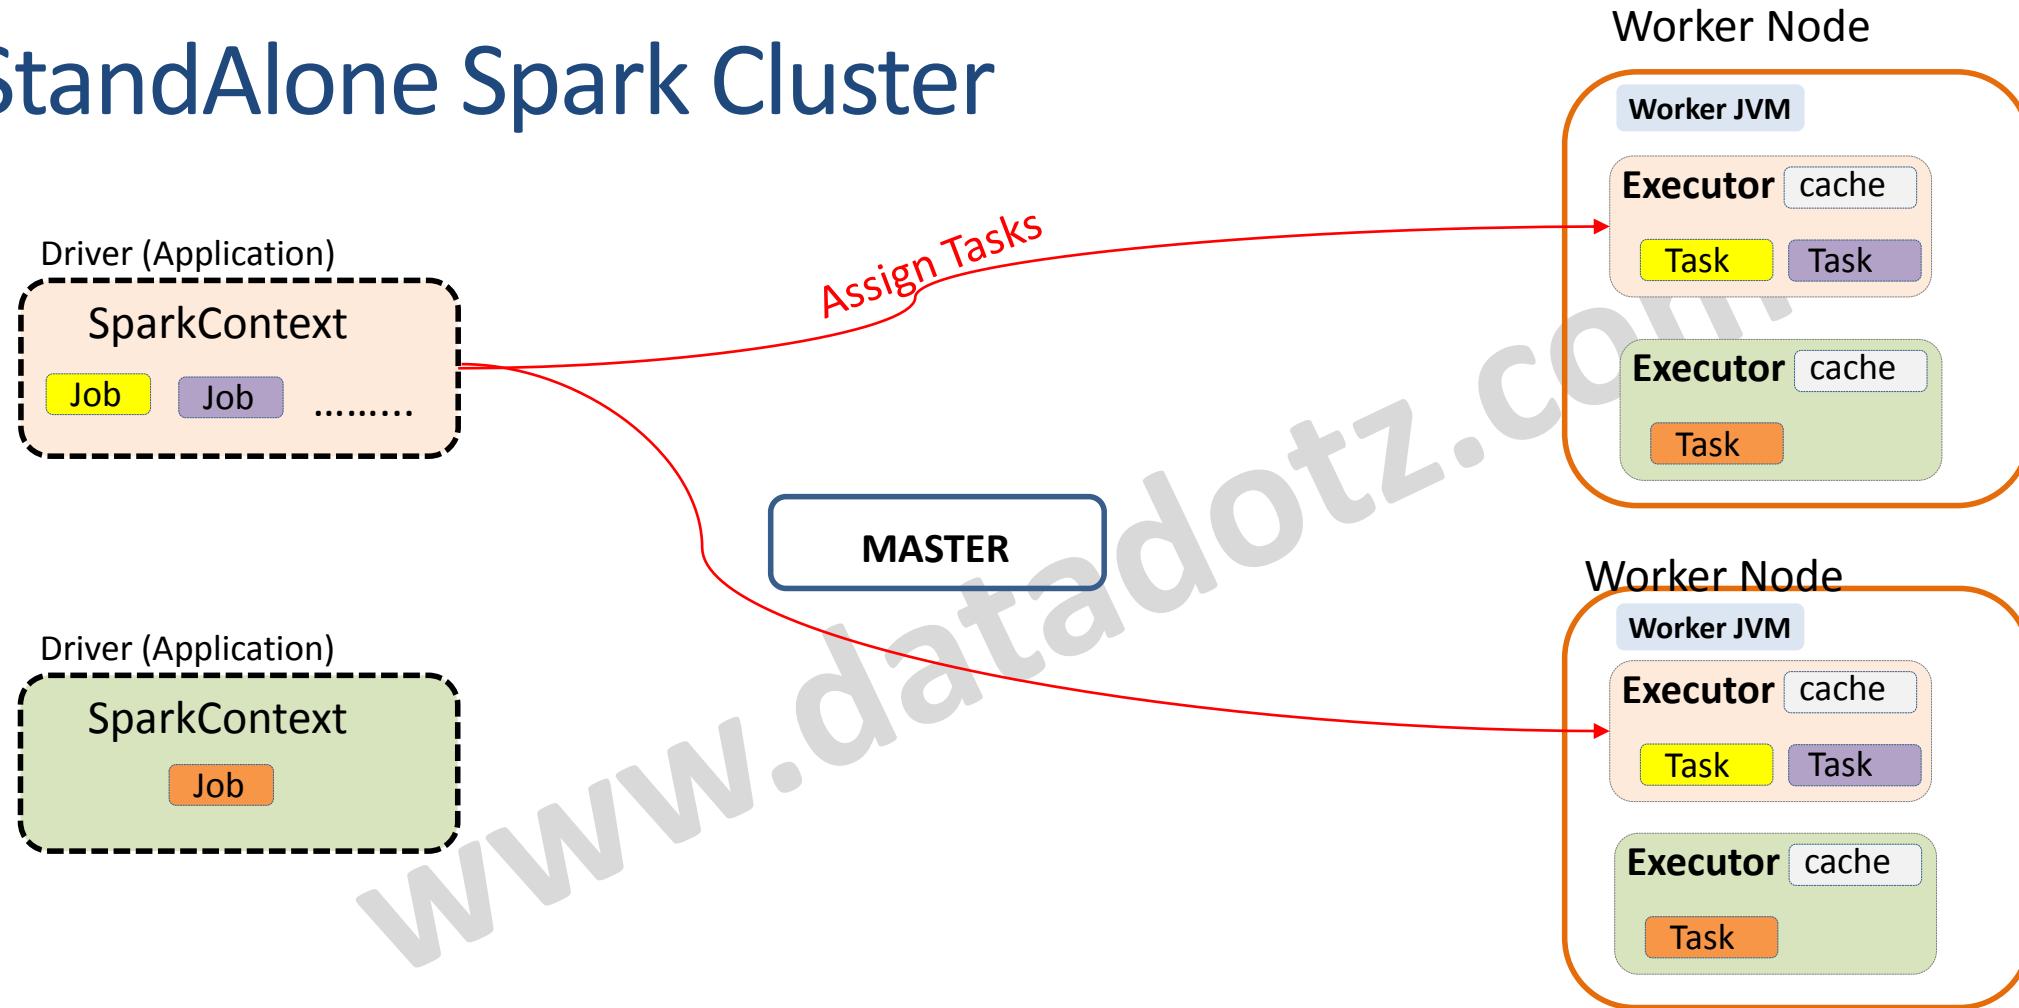
- Application = driver + its executors(JVM)
  - Will have only one SparkContext
- Central Coordination for the spark application in Driver
- Responsible for RDD logical & physical plan creation
  - Assigns Tasks to executors

# Executors



- Launched by Worker (on cmd from Master) in Standalone
- Runs Tasks assigned by Driver
  - Tasks correspond to Stages
  - Tasks work on their respective RDD Partitions
- Cache RDDs whenever needed


# StandAlone Spark Cluster



Scheduling – Covered Later\*\*

# Memory Allocation in Executor

■ RDD Storage ■ Shuffle Storage ■ User Programs



## RDD Storage:

- Use `Spark.storage.memoryFraction` to limit use JVM heap
- Used for `.persist()` or `.cache()`

## Shuffle Storage:

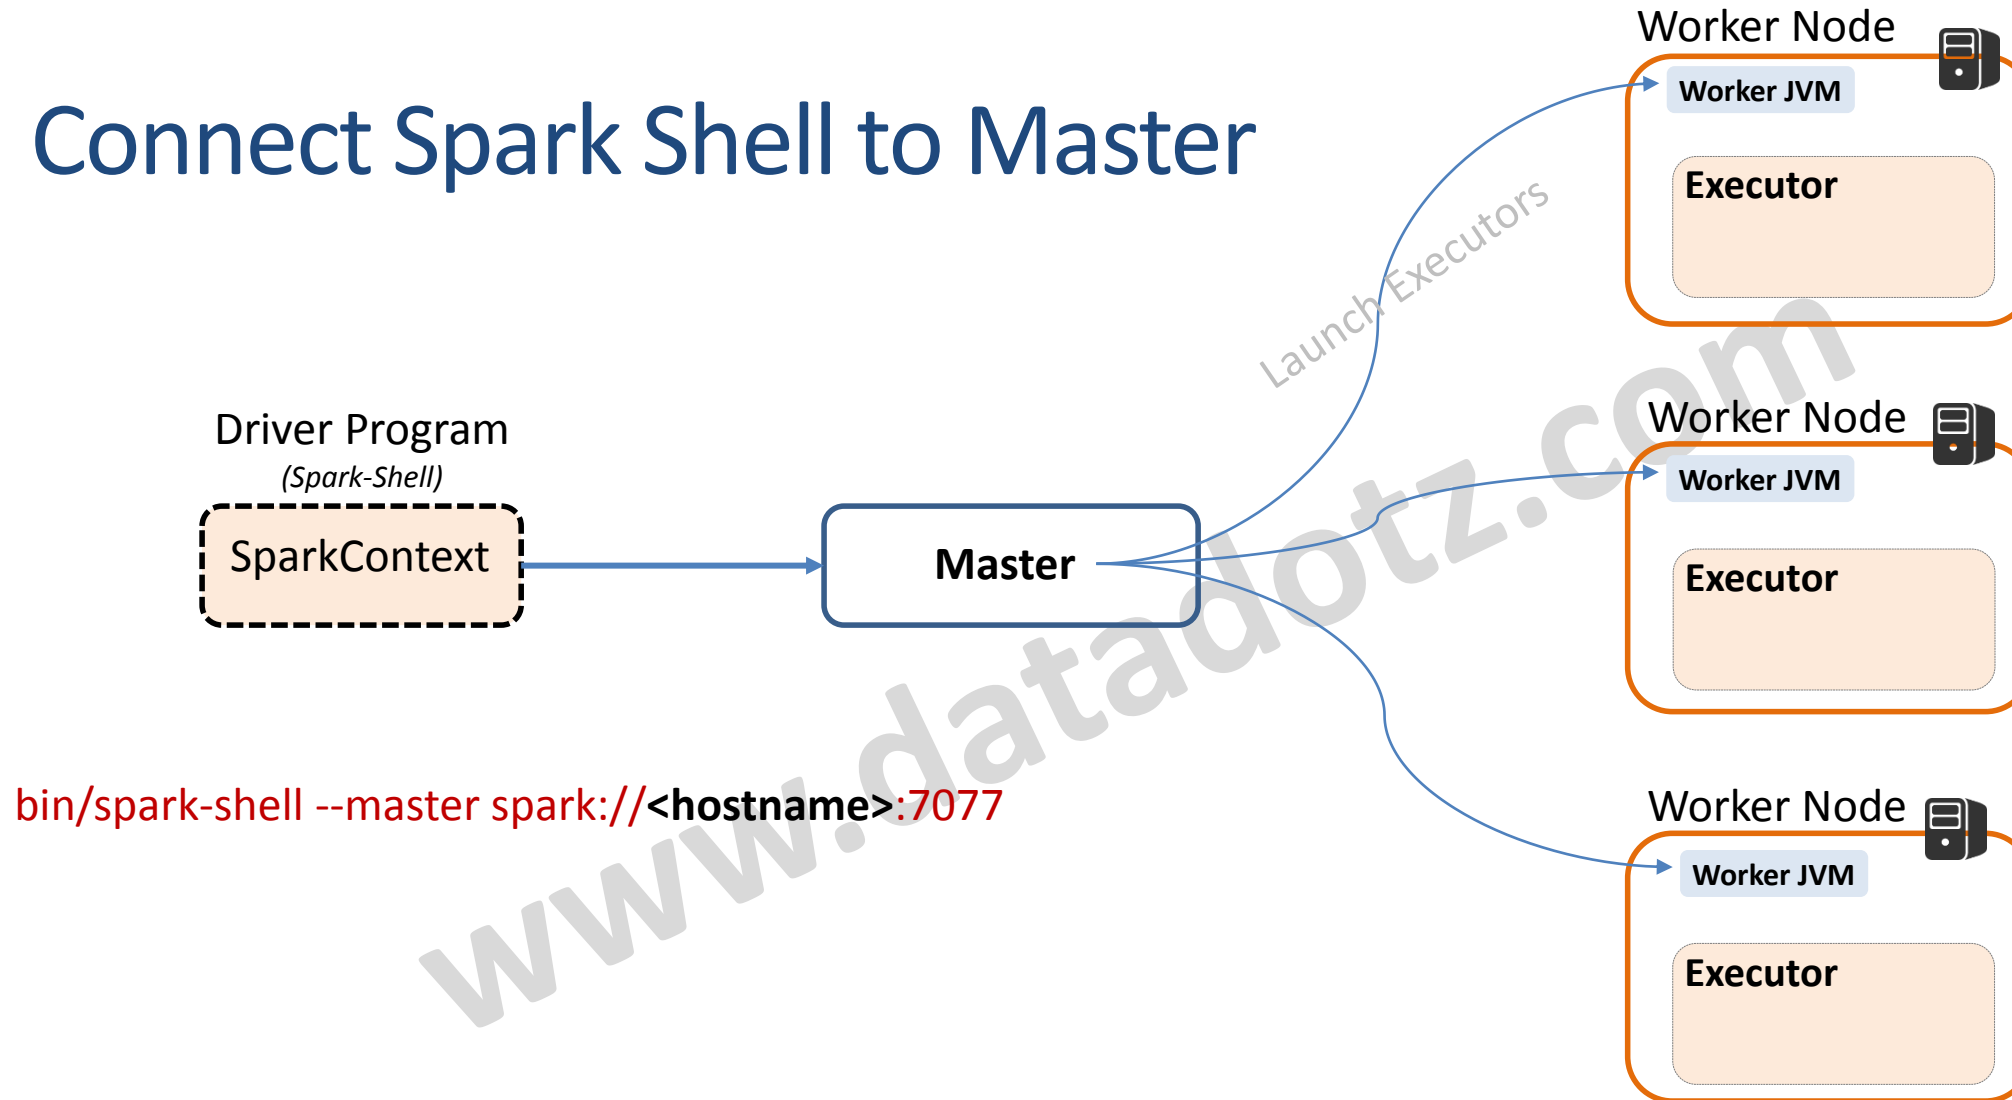
- Intermediate buffer for shuffle output data

## User Programs:

- JVM memory for executing user code.




# Connect Spark Shell to Master



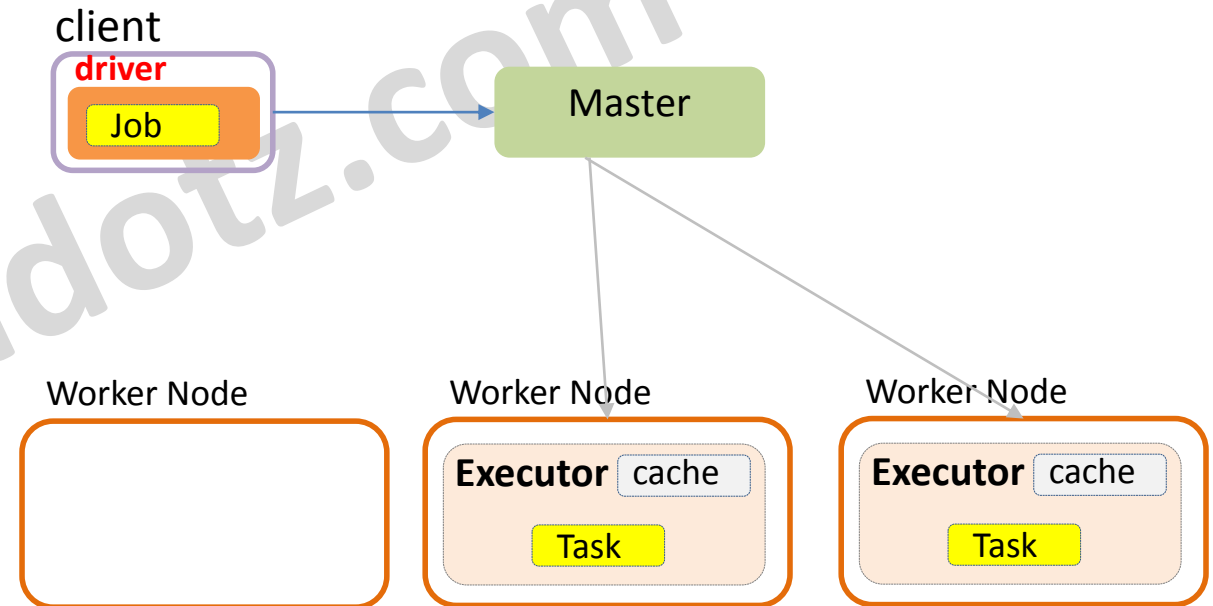
# Deploy-mode

Cluster mode



Fire & Forget

Client mode



# Spark Submit

**bin/spark-submit [options] <app jar> <app jar arguments>**

## Common Options

- master
- class
- deploy-mode
- name
- jars
- conf
- properties-file
- executor-memory
- total-executor-cores

# Types of Application

- Long Lived / Shared Application
  - SparkSQL JDBC Applications - ThriftServer
  - Spark Streaming
- Short Lived Applications
  - StandAlone Applications
  - Interactive Shell Sessions

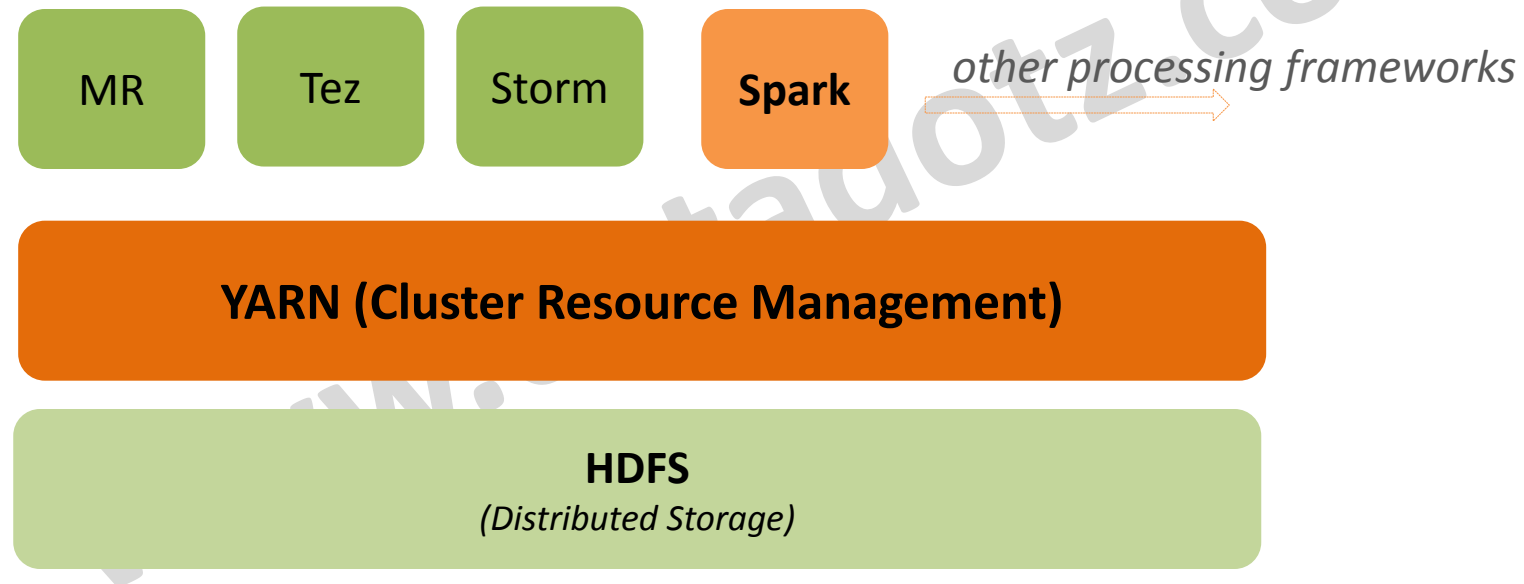
# Spark and Hadoop

- Supports below inputs/outputs from Hadoop
  - Text Files
  - Sequence Files
  - Avro
  - Parquet
  - Other InputFormat

**Spark**  
StandAlone

**HDFS**  
(Distributed Storage)

# Hadoop 2. X



YARN - Yet Another Resource Manager  
MR - MapReduce

# Summary of Cluster Managers


	Spark StandAlone (client)	Spark StandAlone (Cluster)	Yarn Client	YARN cluster
<b>Driver runs in</b>	Client	As JVM in Worker	Client	Application Master
<b>Resources requested by</b>	Client	Driver JVM in Worker	Application Master	Application Master
<b>Executors started by</b>	Worker	Worker	Node Manager	Node Manager
<b>Daemons</b>	Master & Slave	Master & Slave	ResourceManager Node Manager	ResourceManager NodeManager
<b>Spark Shell</b>	Yes	No	Yes	No

# Hadoop Vendor Support

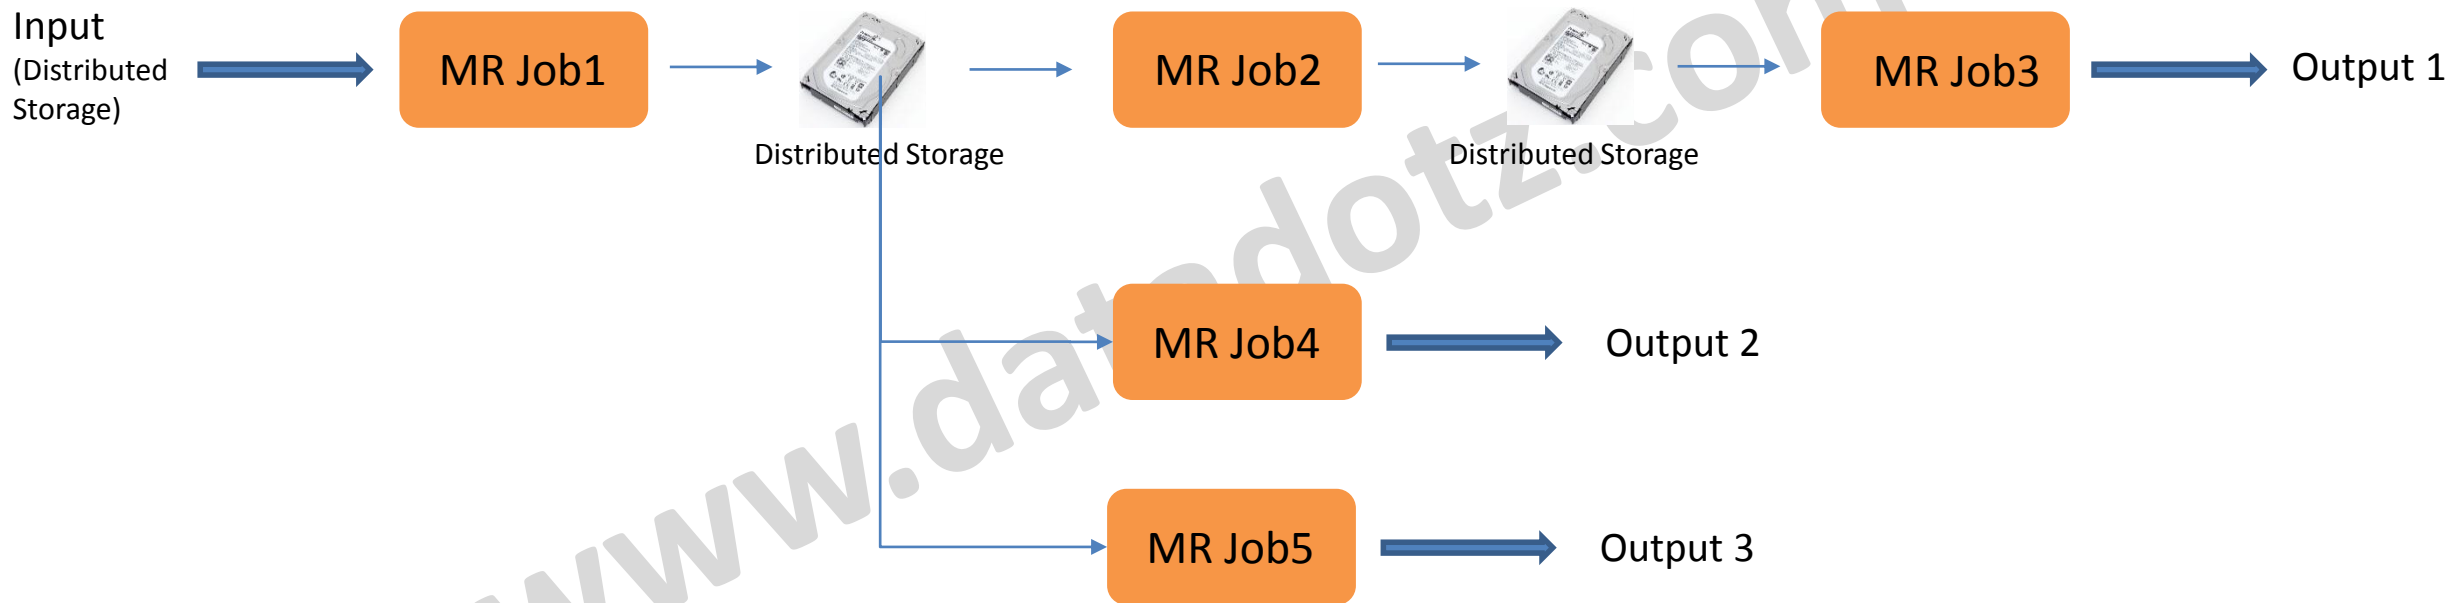




# Lets Understand with an Example



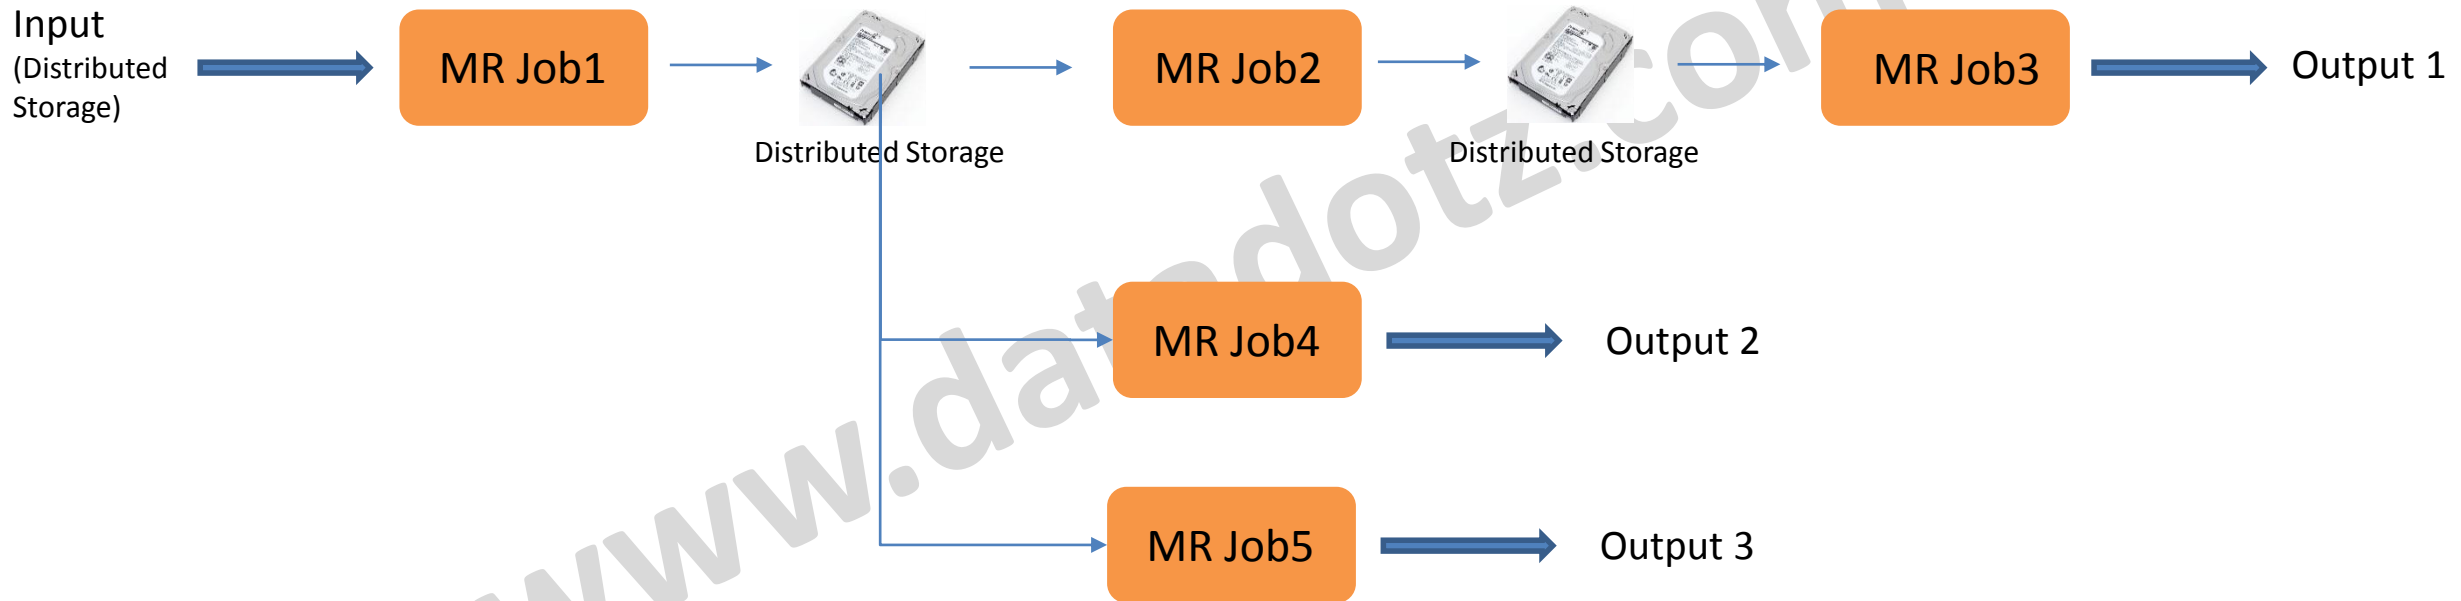
# MapReduce



- Inefficient for Data Reuse in Applications
  - Iterative / MultiStep Applications
  - Interactive Applications

Reference: [http://www.cs.berkeley.edu/~matei/papers/2012/nsdi\\_spark.pdf](http://www.cs.berkeley.edu/~matei/papers/2012/nsdi_spark.pdf) - Research from Matei Zaharia and Others

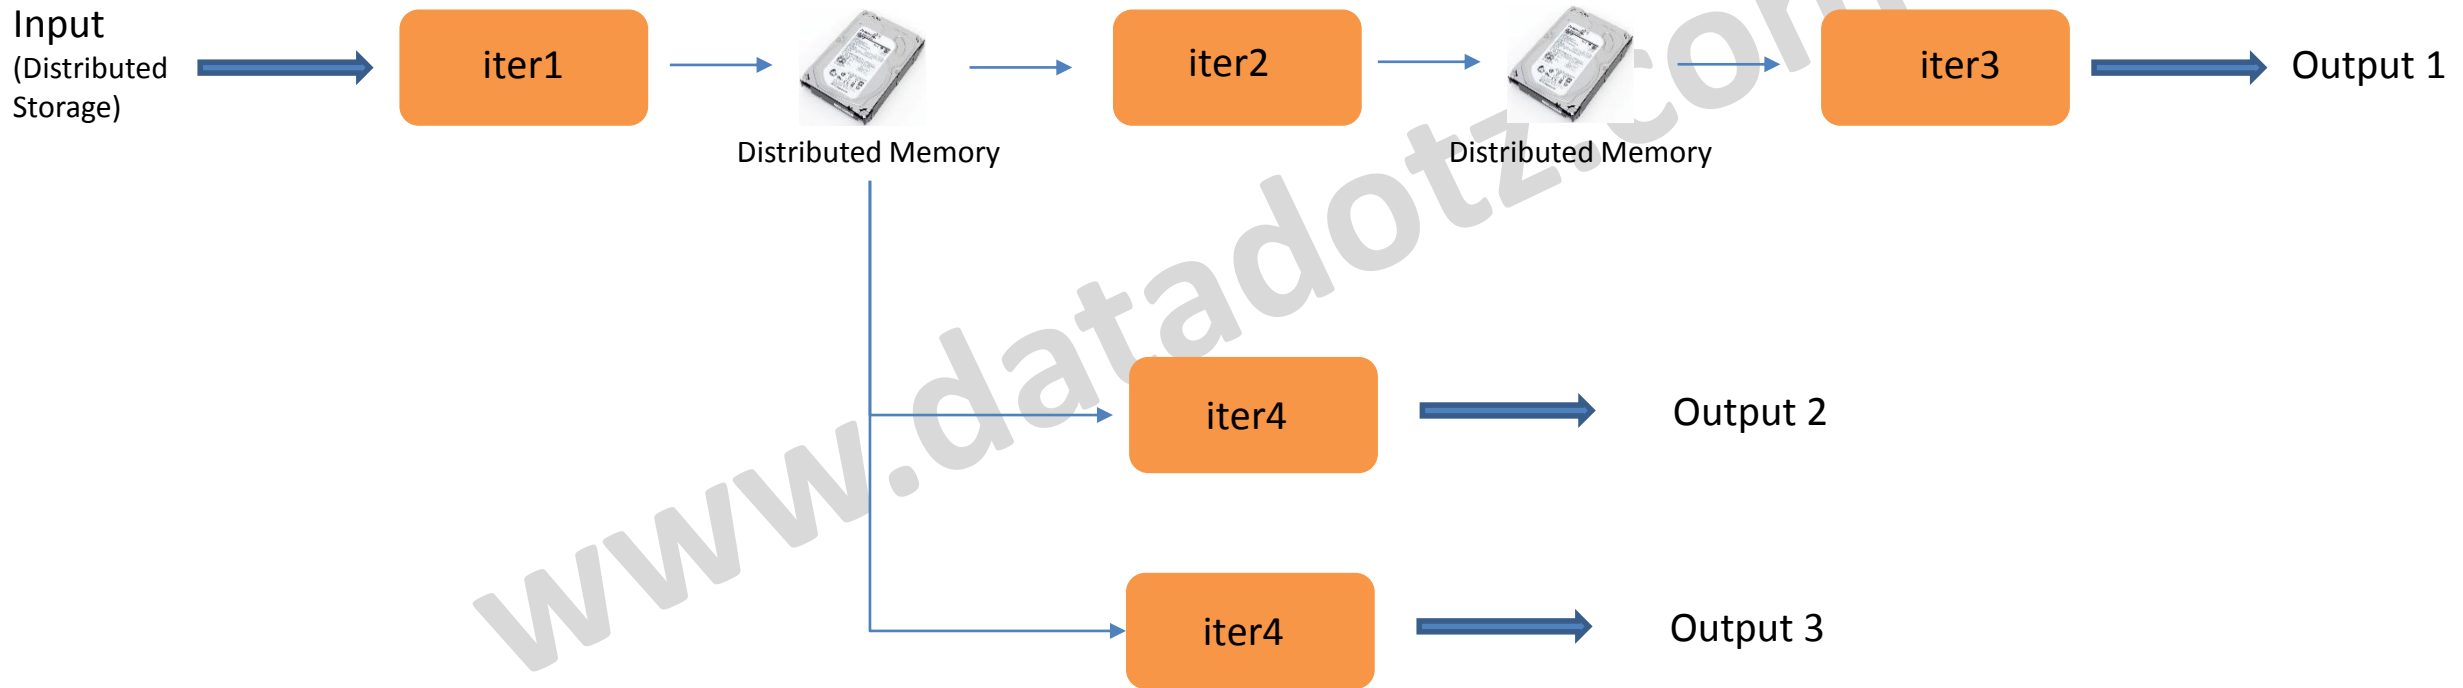
# MapReduce



- Inefficient for Data Reuse in Applications
  - Iterative / MultiStep Applications
  - Interactive Applications

Reference: [http://www.cs.berkeley.edu/~matei/papers/2012/nsdi\\_spark.pdf](http://www.cs.berkeley.edu/~matei/papers/2012/nsdi_spark.pdf) - Research from Matei Zaharia and Others

# Spark



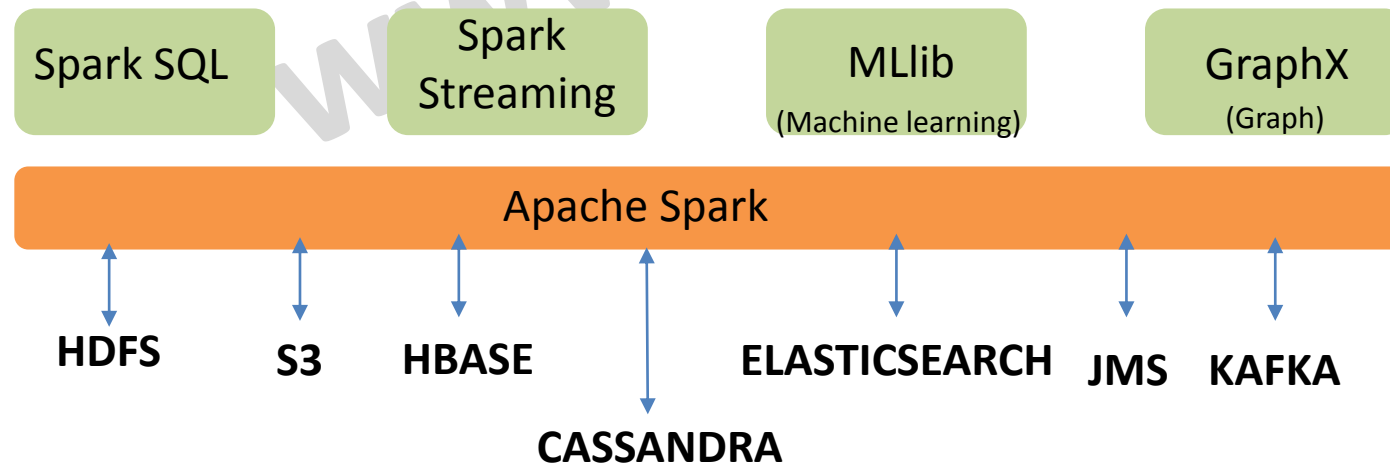
Reference: [http://www.cs.berkeley.edu/~matei/papers/2012/nsdi\\_spark.pdf](http://www.cs.berkeley.edu/~matei/papers/2012/nsdi_spark.pdf) - Research from Matei Zaharia and Others

# Spark is faster

- Caching – Intermediate Data - In Memory
  - Iterative or MultiStep Algorithms
  - Workflows are faster due to caching
- Shuffling
  - Low cost
  - Hash, Sort, Tungsten-Sort
- Startup time
  - Tasks are Threads in Spark whereas Tasks are JVMs in MR

# Other factors

- Multiple operators
  - MR – Map, Reduce
  - Spark – map, reduce, join, filter, cogroup, sort..etc
  - **Ease of Programming – RAD(Rapid Application Development)**
- Execution Engines
  - MR – Batch
  - Spark – Batch, streaming, Interactive



**ThAnK yOu**